

# ECE 485/585

# Microprocessor System Design

Prof. Mark G. Faust

Maseeh College of Engineering  
and Computer Science



# Virtual Memory

ECE 485/585

Mark G. Faust

# What is Virtual Memory

- Automatic address translation
  - Decouples program's addresses from physical location
  - Provides for program address space bigger than physical memory
  - Allows expanding program's address space without re-allocating existing memory
  - Provides protection from other tasks/processes
- Components of virtual memory
  - Physical memory partitioned into pages
  - Swap device (typically disk) holds pages not resident in physical memory
  - Address translation from program (virtual) address to physical address
    - Page tables
      - Provide translation from virtual to physical address (if resident)
      - Indicate where page can be found on swap device (backing store) if not resident
    - Translation lookaside buffer (TLB) is a cache of translations
  - Hardware/Software cooperation
    - TLB in hardware
    - Page table look up in H/W or S/W
    - Page table management in S/W

# Origins of Virtual Memory

Early days of computing (pre-1960)

Single task batch systems: dedicated to one user

Unused memory space is wasted

Time spent waiting for I/O is wasted (CPU is idle)

Multi-task batch systems:

Run another task while one is waiting for I/O

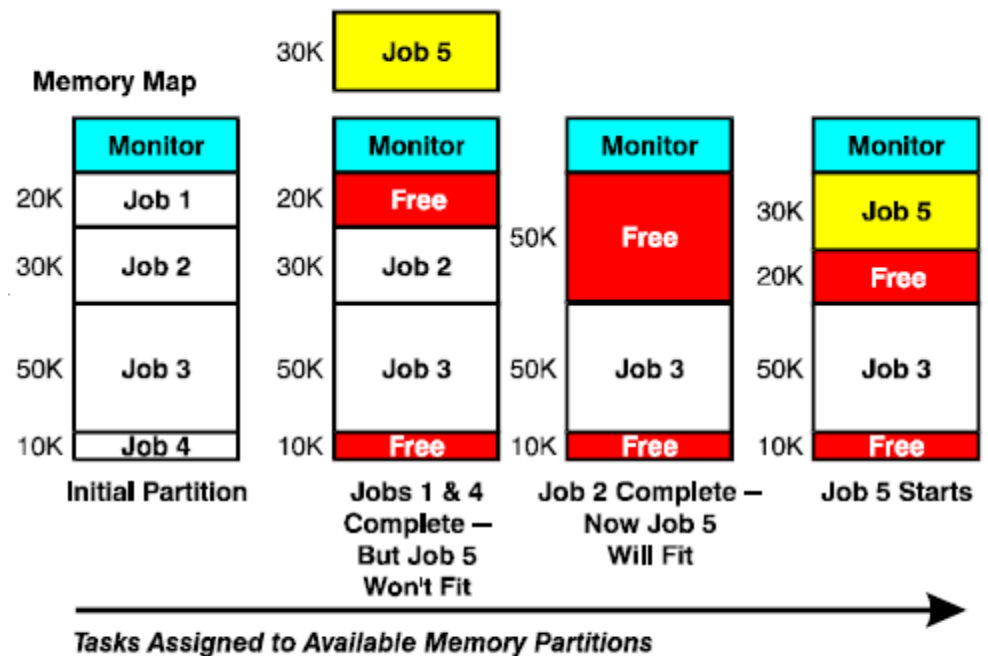
Multiple programs resident in memory

New program needs to await free space

must be contiguous

can result in waits even when space available

Memory Map



# Origins of Virtual Memory

Relocation registers (IBM S/360 c. 1964)

Base register is added to all memory references

Rudimentary protection

- No addresses less than base register

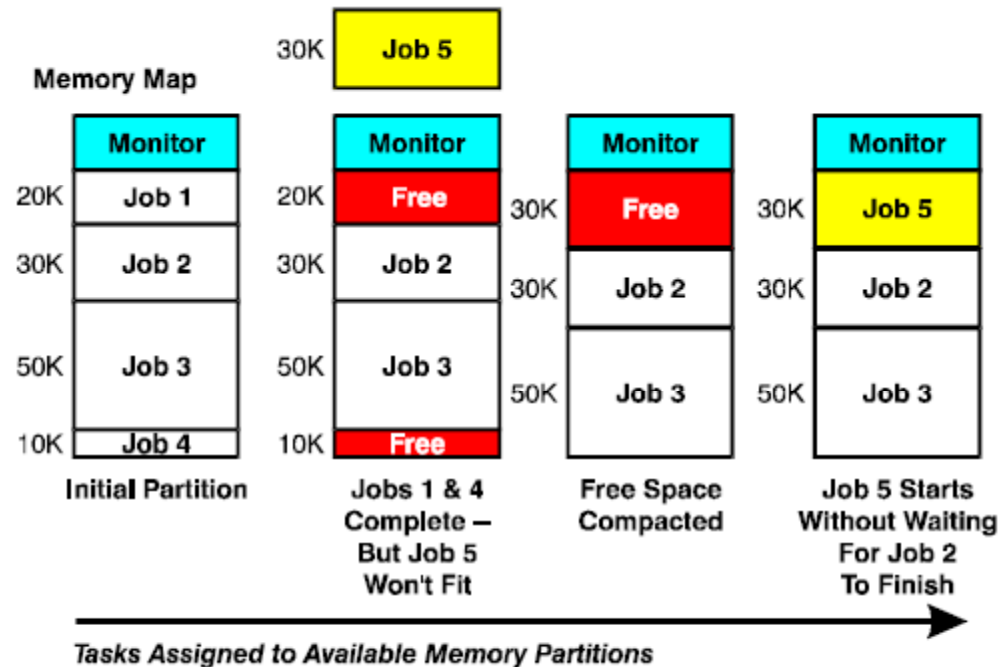
- Exception if Base + Address exceeds Bound

Programs can be moved in physical memory

- if we adjust base register

Still consume time by copying memory image

Doesn't allow programs > than physical memory




# Origins of Virtual Memory


## Use of overlays

Programs partitioned into segments which were overlayed in physical memory (to/from disk) as needed (entire program didn't need to reside in physical memory at the same time)

### User program



### Physical memory




Done manually by programmer

Tedious and very difficult to do optimally to minimize swapping of overlays

# Virtual Addressing

Early 1960s...

Physical addressing

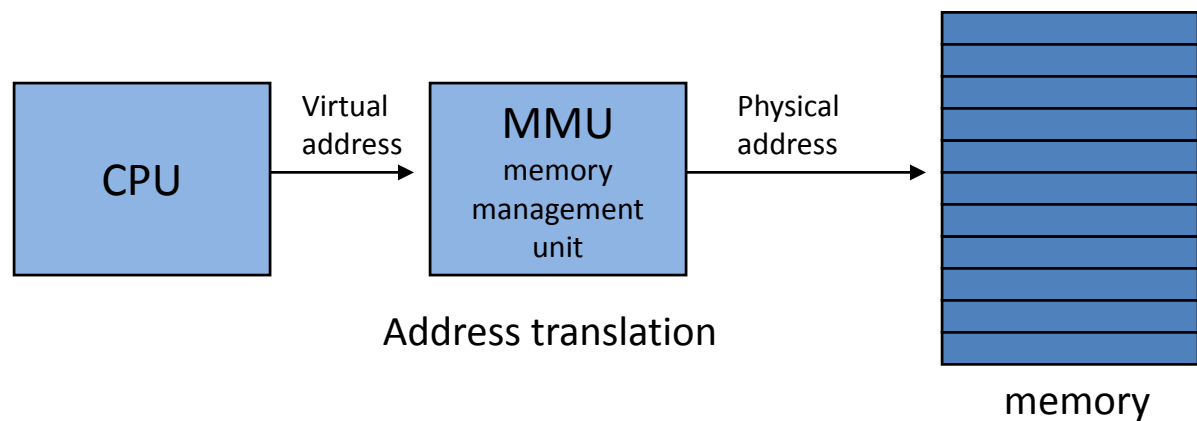


Mapping from virtual to physical addresses


Virtual address space could be larger than physical address space

Each process could have its own virtual address space

Virtual addressing



# Segmentation



Address is (Segment, Offset) pair


Each segment has own contiguous address space

Facilitates relocating code

Segment can be relocated anywhere in physical memory  
position independent code!


Offsets in segment maintained

Physical address is Base Address of Segment + Offset



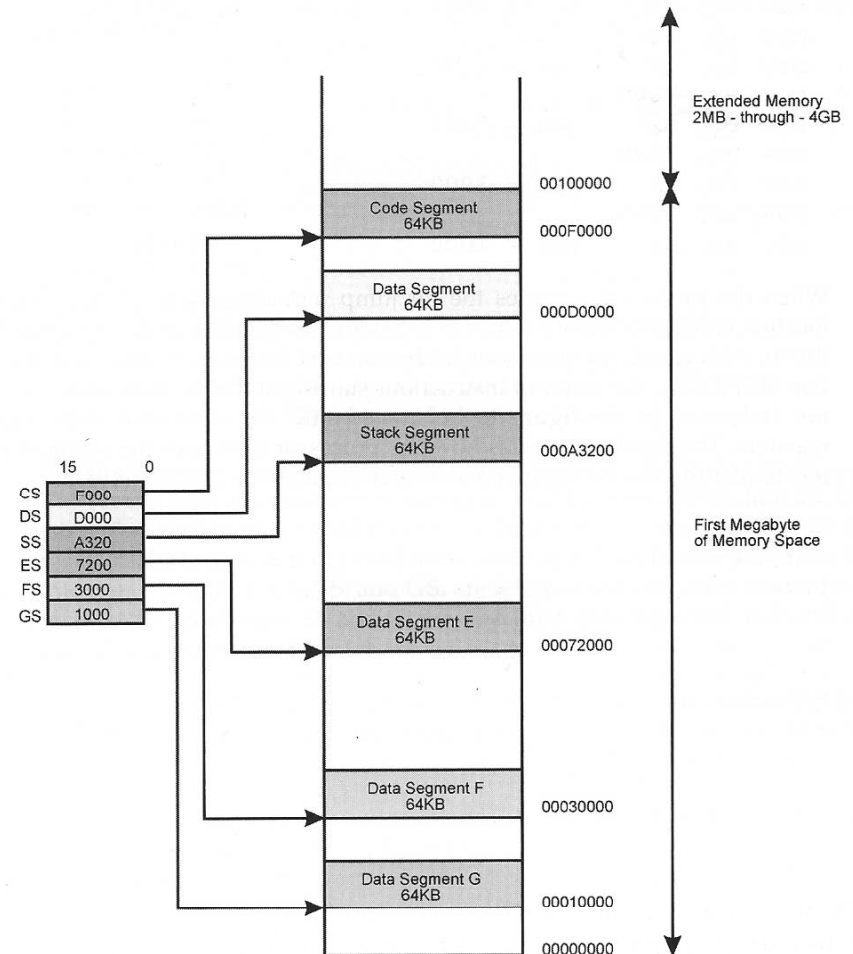


# Segmentation




# Segmentation: Intel

- Code and data references
  - Segment:Offset e.g. CS:IP
- Segment registers
  - Contain base address of segment
    - Code Segment (CS)
    - Default Data Segment (DS)
    - Stack Segment (SS)
    - Additional Data Segments
- Simplest case (e.g. Real Mode)
  - Linear Address =
    - Segment Base (shifted) + Offset



# Real Mode Segmentation

- Default Data Segment Register (DS)
  - Contents shifted left 4 to form base address
  - Creates effective 20 bit base address
- Offset (actual operand)
  - Added to shifted base address to form linear address
- Same process occurs for
  - Normal instruction execution
    - Code Segment (CS)
  - Stack operations
    - Stack segment (SS)
- Explicit segment references allowed
  - “Segment overrides”
  - `MOV BH, ES:[0002]`



# Segment-Level Protection

- Segmentation evolved to provide additional support
  - Protect OS and other processes from each other
    - Malicious process
    - Buggy process
  - Aid to debugging
    - Process memory image not corrupted
    - Debugger resident and not corrupted
- Every memory reference is checked
  - Type Check (e.g. don't write to code segment)
  - Limit Check (e.g. offset within segment's range)
  - Restriction of addressable domain (through segments)
  - Restriction of procedure entry points (can't jump/branch to random places in other segments – e.g. system code!)
  - Restriction of instruction set (only OS can manipulate segment registers, etc)

# Segment-Level Protection

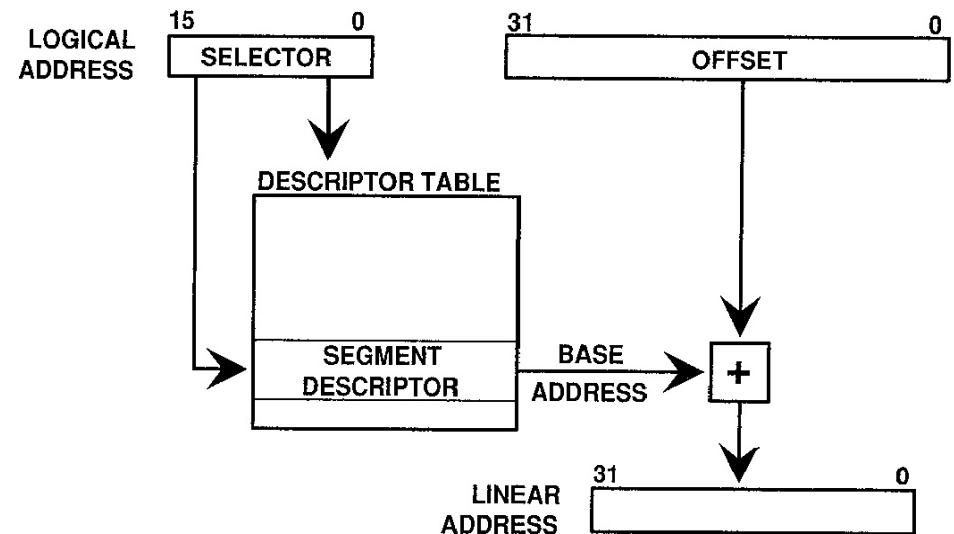
- Restriction of instruction set
  - Reserved to Operating System (Level 0)
    - HLT Halt
    - INVD Invalidate the cache
    - INVLPG Invalidate TLB Entry
    - LGDT Load GDT register
    - LIDT Load IDT register
    - LLDT Load LDT register
    - LMSW Load Machine Status Word
    - LTR Load Task Register
    - MOV <CRn> Move to Control Register
    - MOV <DRn> Move to Debug Register
    - WBINVD Write Back and Invalidate Cache

# Segment-Level Protection

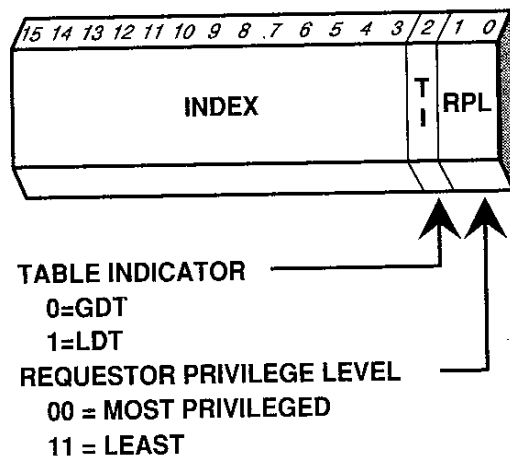
- Restriction of instruction set
  - IOPL (I/O protection level)
  - “Sensitive” instructions reserved to device drivers
    - IN                      Input
    - INS                    Input String
    - OUT                    Output
    - OUTS                  Output String
    - CLI                    Clear Interrupt-Enable Flag
    - STI                    Set Interrupt-Enable Flag
- Other checks
  - Stack on IRET and RET
    - Don’t allow user process to set-up bogus return IP on stack

# Protected Mode Segmentation

- Segment registers are used as selectors instead of actual base address
  - Index into either GDT or LDT
    - Global or Local Descriptor Table
    - Single GDT
    - 0, 1, or many LDTs
- Descriptors contain segment info
  - Base address
  - Limit
  - Privileges
- Hardware performs bounds checking
  - Data and Code Segments
- Hardware performs access checking
  - Code Segments




# Protected Mode Segmentation



Entry 0 of GDT unused? Why?

New CS register ("selector") loaded upon CALL or long JMP

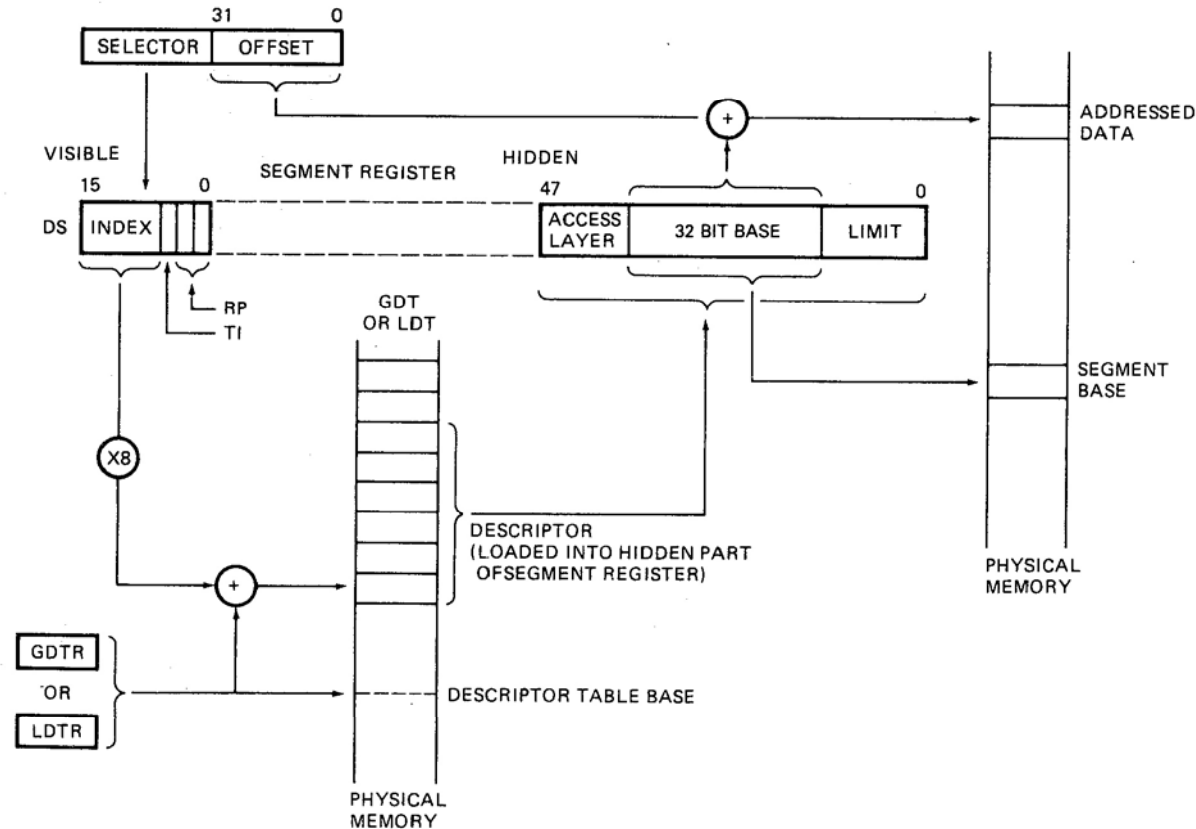
For efficiency, contents of descriptors for current segments as cached in hidden parts of registers



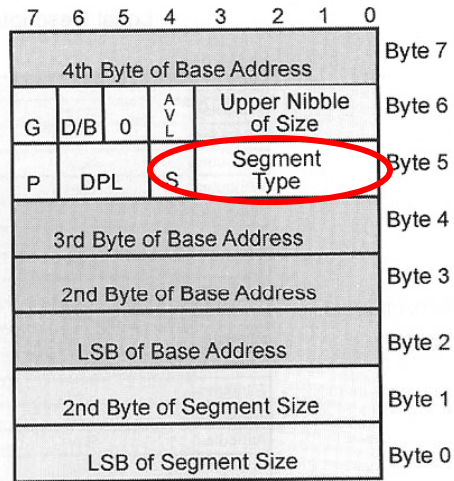


# Descriptors are cached

- Descriptors are cached in hidden portions of the segment registers
- Saves memory accesses for each reference!



# Segment Descriptors

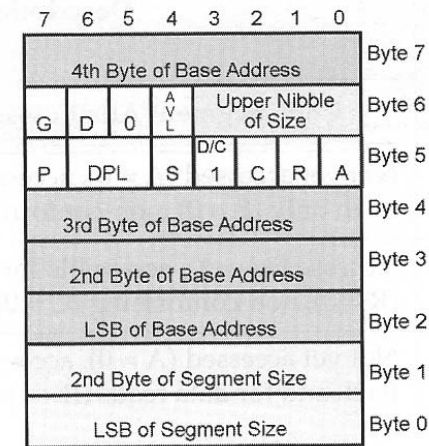


Code  
Data  
Call Gate

G Bit	Granularity bit defines meaning of limit value: 0 = length of segment in bytes. 1 = length of segment in pages.
Type Field	Defines a segment as data or code, read-only or read/write, expansion direction and whether it is has been accessed or not.
D/B Bit	In a code segment, D bit indicates default operand and address size. In data segment, B bit defines SP size and upper bound of expand-down stack.
AVL Bit	Available for use by system software.
P Bit	Segment Present bit.
DPL Field	Descriptor Privilege Level.
S Bit	System bit. When 0, indicates system segment.
Size Field	Specifies the size of the segment (also see the G bit).
Base address	Specifies the 32-bit base linear address of the segment in the 4GB linear address space.

# Protected Mode Segmentation

- Code Segment Descriptors
  - 32-bit base address
  - 20-bit segment size (limit)
    - G bit (granularity) indicates whether limit is
      - Bytes
      - 4K pages ( $4K = 2^{12}$ ) therefore  $2^{32} = 4GB$
- Types of code segments
  - Conforming
    - Used for shared libraries
    - Segment conforms to privilege of caller
  - Non-Conforming



**G Bit** Granularity bit defines meaning of limit value.  
 0 = length of segment in bytes.  
 1 = length of segment in 4KB pages.

**D Bit** In code segment, Default bit defines default size of operands and effective addresses. 0 = 16-bit, 1 = 32-bit.

**AVL Bit** Available for use by system software

**P Bit** Segment Present bit (must be 1 if the code segment is present in memory).

**DPL Field** Descriptor Privilege Level (0-3)

**S Bit** System bit. When 0, indicates system segment. Must be 1 in a code segment descriptor.

**D/C** This could be called the Data/Code bit. A 0 indicates a data segment and a 1 indicates a code segment.

**C Bit** Conforming bit. Set to 1 if code segment is conforming. See text for a detailed description.

**R Bit** Readable bit. A 0 indicates an execute-only segment, while a 1 indicates the segment may be read from by both the prefetcher and for data accesses.

**A Bit** Accessed bit. Set to 1 by the processor when a code segment is accessed.

# Protected Mode Segmentation

- Data Segment Descriptors

	7	6	5	4	3	2	1	0
Byte 7	4th Byte of Base Address							
Byte 6	G	B	0	A V L	Upper Nibble of Size			
Byte 5	P	DPL		S	C/D 0	E	W	A
Byte 4	3rd Byte of Base Address							
Byte 3	2nd Byte of Base Address							
Byte 2	LSB of Base Address							
Byte 1	2nd Byte of Segment Size							
Byte 0	LSB of Segment Size							

**G Bit** Granularity bit defines meaning of limit value:  
0 = length of segment in bytes.  
1 = length of segment in pages.

**B Bit** In data segment, Big bit defines SP size and upper bound of expand-down stack.

**AVL Bit** Available for use by system software

**P Bit** Segment Present bit (must be 1 if the data segment is present in memory).

**DPL Field** Descriptor Privilege Level

**S Bit** System bit. When 0, indicates system segment. Must be 1 in a data segment descriptor.

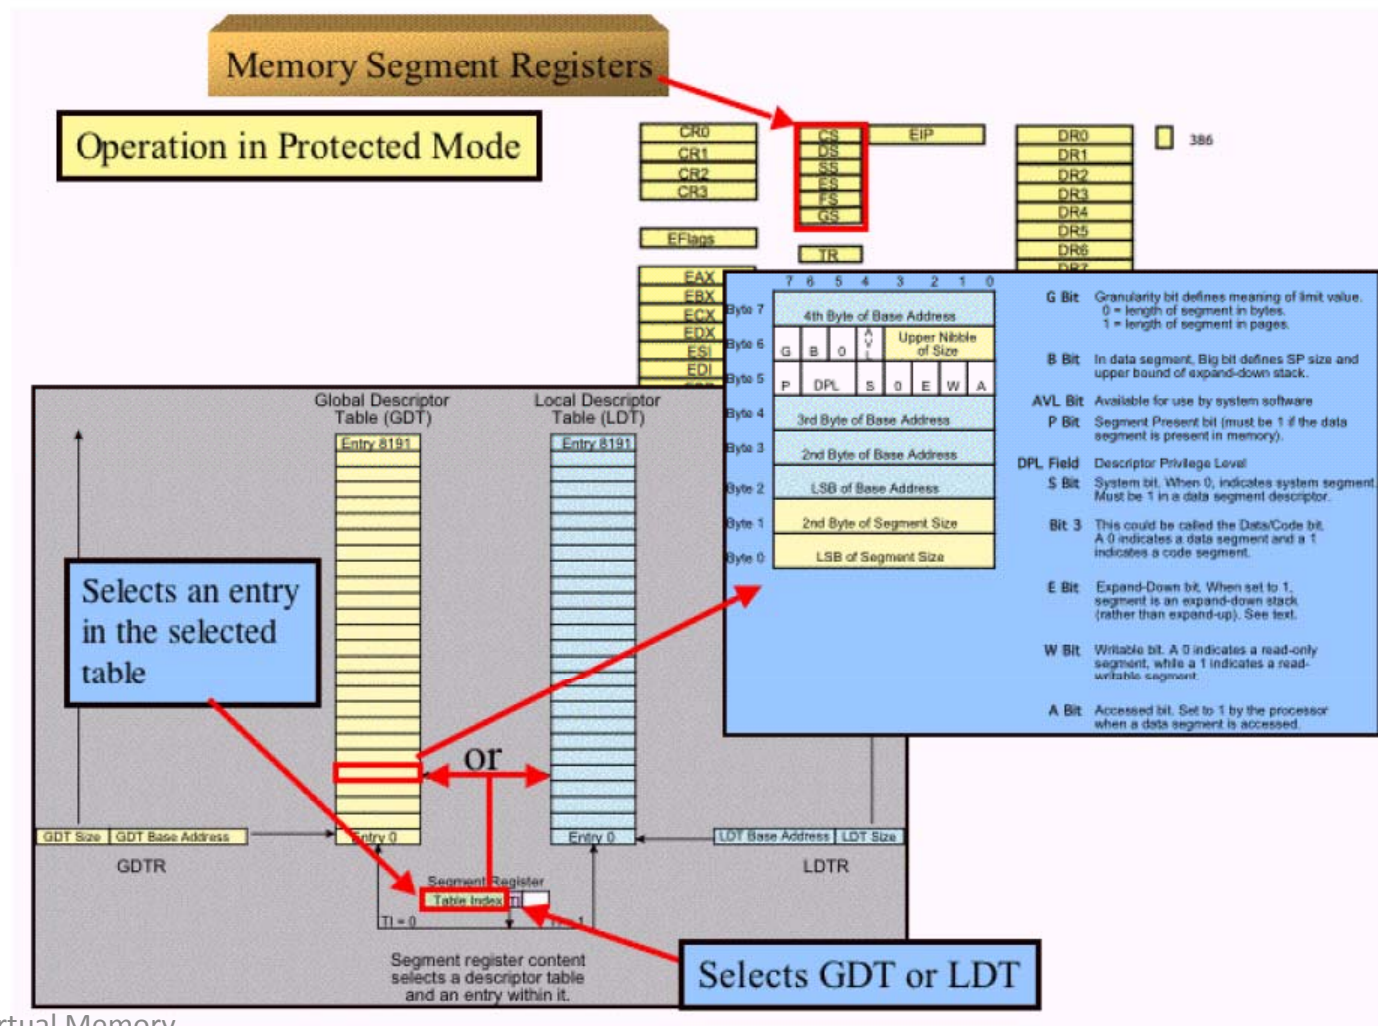
**Bit 3** This could be called the Data/Code bit. A 0 indicates a data segment and a 1 indicates a code segment.

**E Bit** Expand-Down bit. When set to 1, segment is an expand-down stack (rather than expand-up). See text.

**W Bit** Writable bit. A 0 indicates a read-only segment, while a 1 indicates a read-writable segment.


**A Bit** Accessed bit. Set to 1 by the processor when a data segment is accessed.

# Protected Mode Segmentation



# Paging

Virtual Address



Segments too large to provide sufficient granularity

- entire segment had to be in memory
- variable size segments made allocation of physical memory difficult

So... paging was created.

Basic idea:

Virtual addresses mapped to physical addresses

Data corresponding to a virtual address may be

- in physical memory
- on disk
- unallocated

But...managing individual addresses too cumbersome and uses too much memory

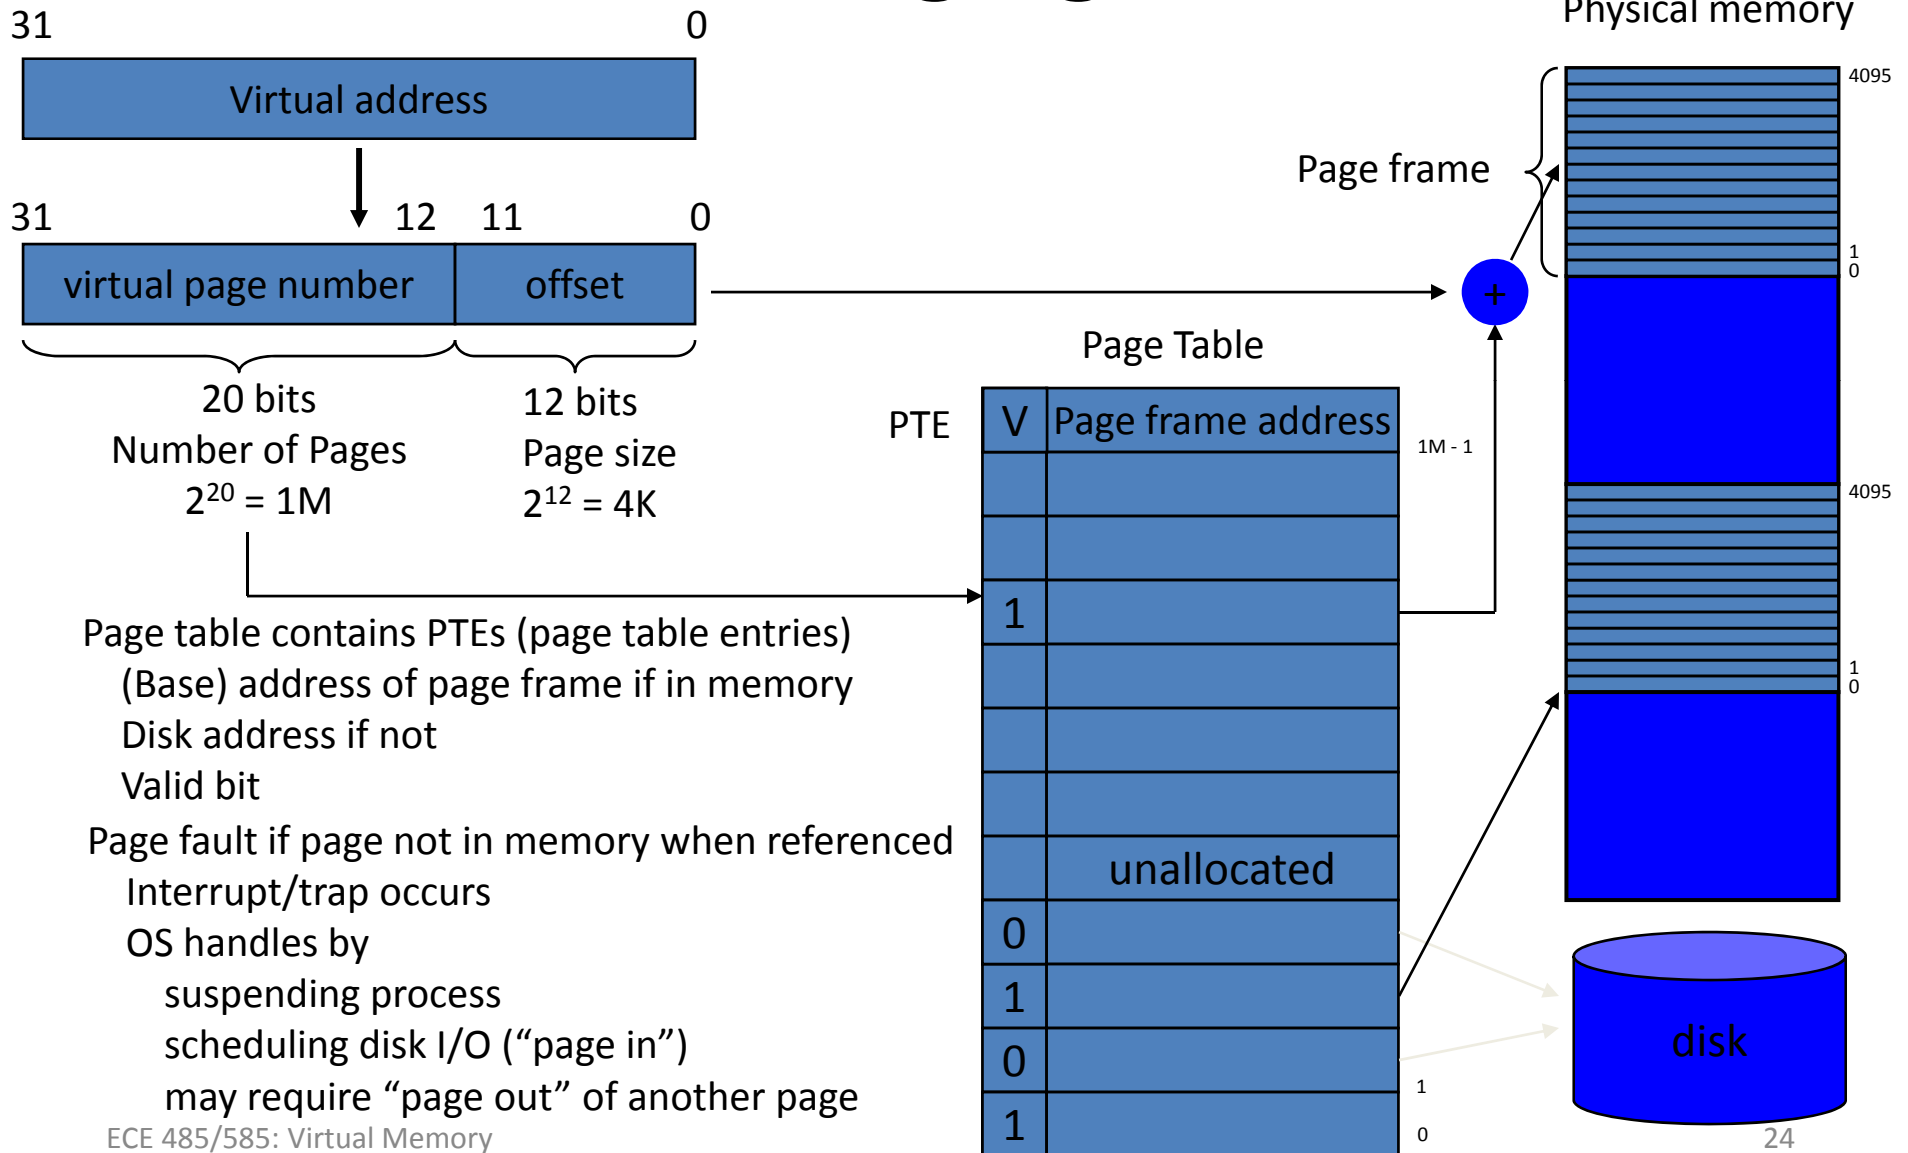
Use fixed-length pages – an entire page is mapped and either resident in memory or on disk

Worked so well, that over time, features of segments incorporated into paging (protection, sharing)

# Benefits of Virtual Memory

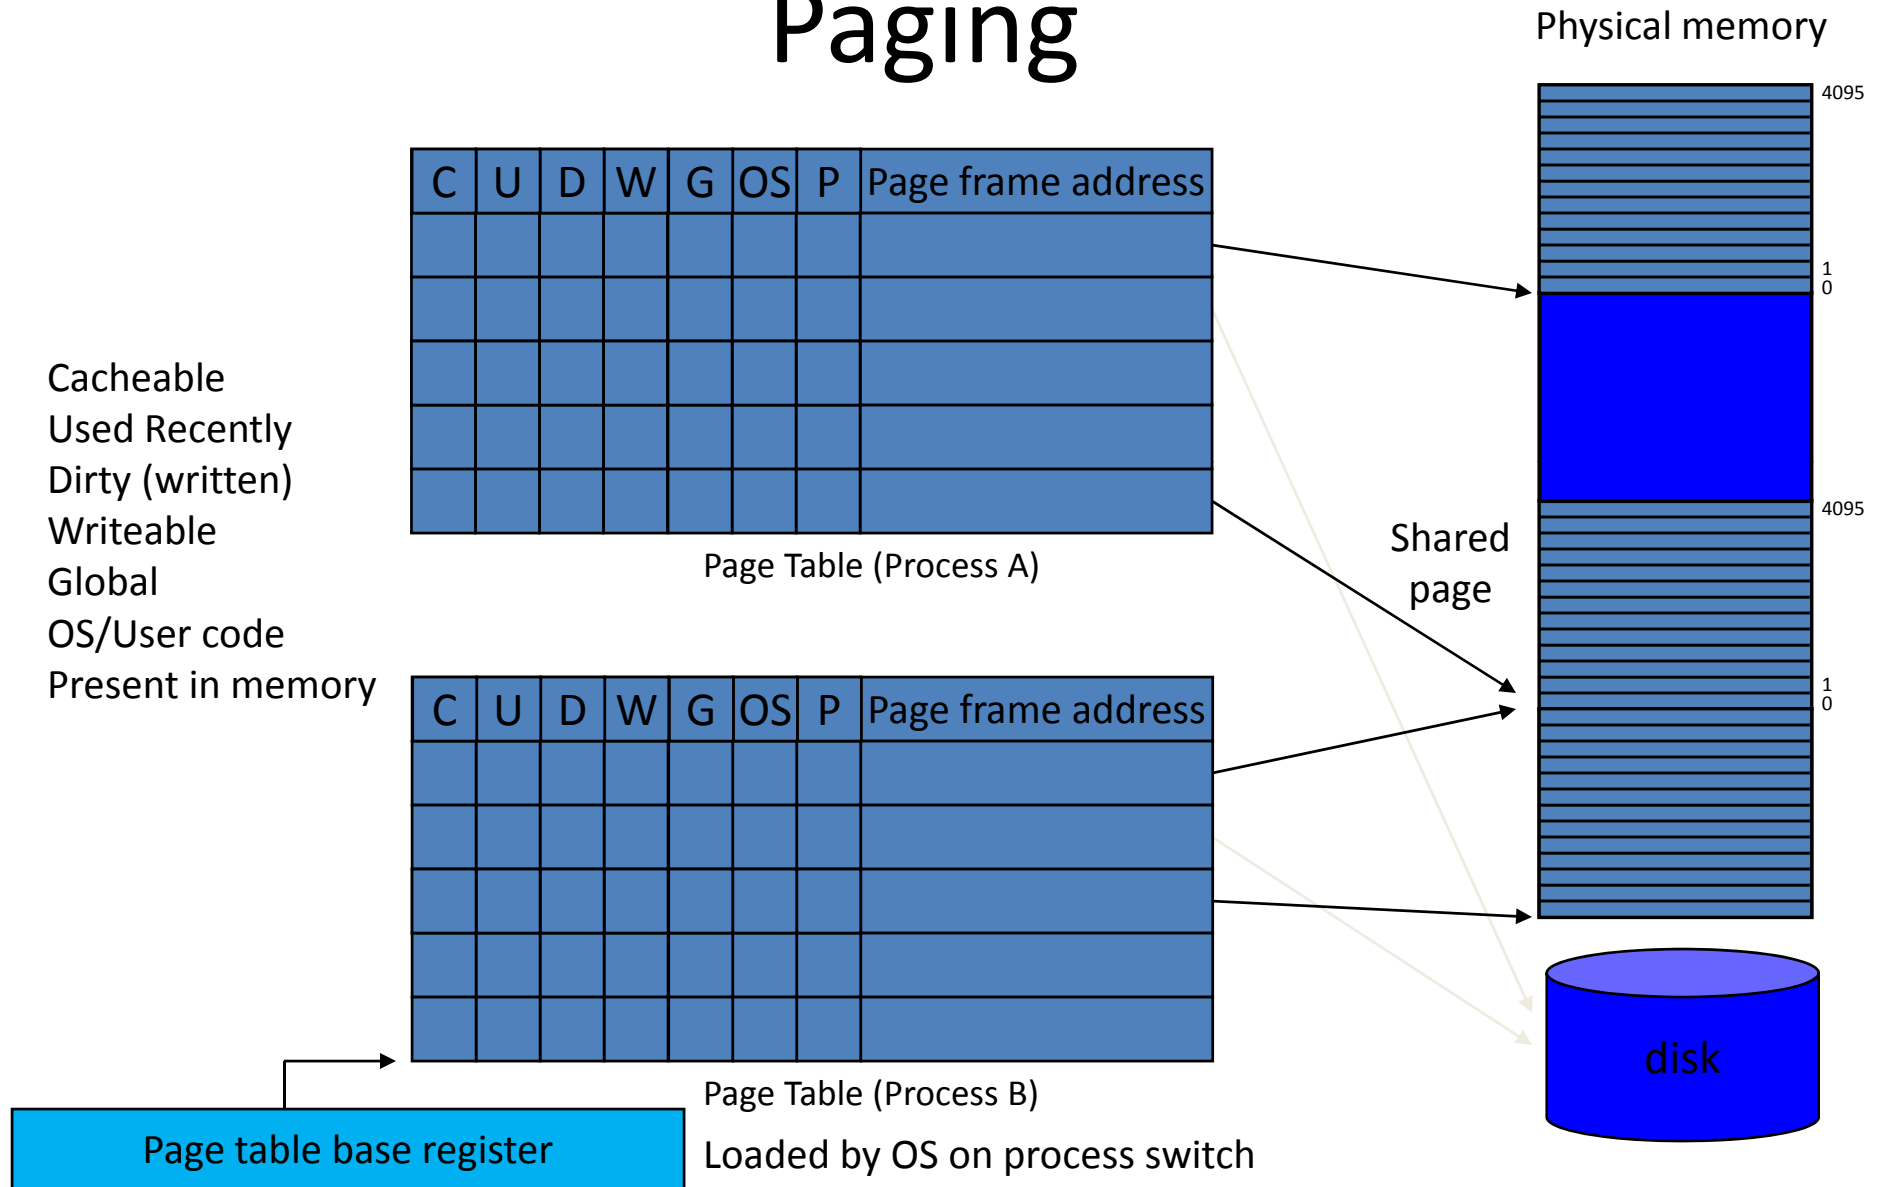
1. Provides every process with large (linear) virtual address space (may be larger than physical memory)
2. Facilitates code sharing (different virtual addresses can map to same physical address, eliminating need to have multiple copies of shared libraries – how many copies of `printf()` in memory?)
3. Facilitates position independent code (data/code can be relocated anywhere in physical memory)
4. Facilitates debugging by providing mechanisms for checking references (e.g. attempts to access unallocated memory, null pointers)
5. Provides protection from intentional or accidental attempts to access or write memory belonging to another user process (or the OS)
6. No need to have entire process memory image loaded to begin execution (“demand paging”)

# Paging





# Paging




# Similarities to Caching


- “Present” bit → “Valid” bit
  - Indicates whether page is present in memory
  - Indicates whether line in cache is valid
- Page Fault → Cache Miss
  - Page not resident in memory, page in from disk
  - Line not in cache, cache line fill from next level memory
- Insufficient memory for new page → Cache collision
  - Insufficient physical memory for page from disk
  - All “ways” in cache occupied for this index
- Need to identify page to be over-written → Victim
  - Page replacement algorithms
  - Cache line replacement algorithms (LRU)
- Need to “page out” pages to make room → Write Back line
- “Dirty” bit to indicate page has been written to → Modified bit
  - No need to write out “clean” pages
  - No need to write back unmodified lines
- Locality and Working Set
  - Avoid “thrashing”

# Types of Page Tables


- Flat
  - Large!
  - Proportional to VA space
  - Too big for physical memory
- Hierarchical
- Inverted (hashed)



# Intel Pentium: Linear (Virtual) to Physical Address Mapping



(a)



# Page Table Lookup


Page Tables can be paged!  
Not first-level “root” or “directory” tables

Each process requires its own page tables. Base address loaded on process switch.

1K x 4K = 4M -- page tables  
4K -- page directory

Page directory can be “sparse” – so don’t require all page tables to be allocated. For example, using 4M virtual memory space requires only 4K (page directory) + 4K (page tables for 1K pages of 4KB each) = 8K.

Multiple memory references/page table lookups for each address (instruction fetch or data reference)



# Translation Lookaside Buffer (TLB)

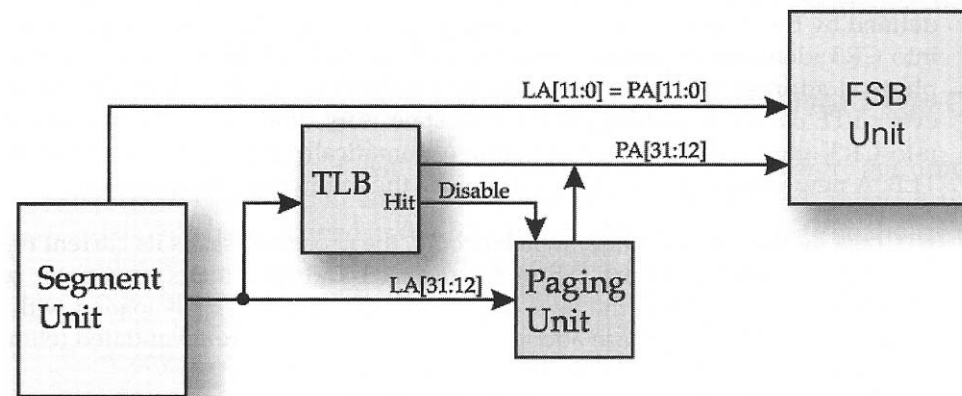
Walking the page table would require multiple memory accesses for each memory reference.

TLB is a cache of most recently used virtual  $\rightarrow$  physical address mappings. Done in parallel with access via paging unit. Similar to... look aside cache!

Single TLB shared by all processes, so must flush the TLB on a context switch (each process has own linear  $\rightarrow$  physical address mapping!) [Intel x86: write to CR3 register flushes TLB] or tag entries with process ID (ASID).

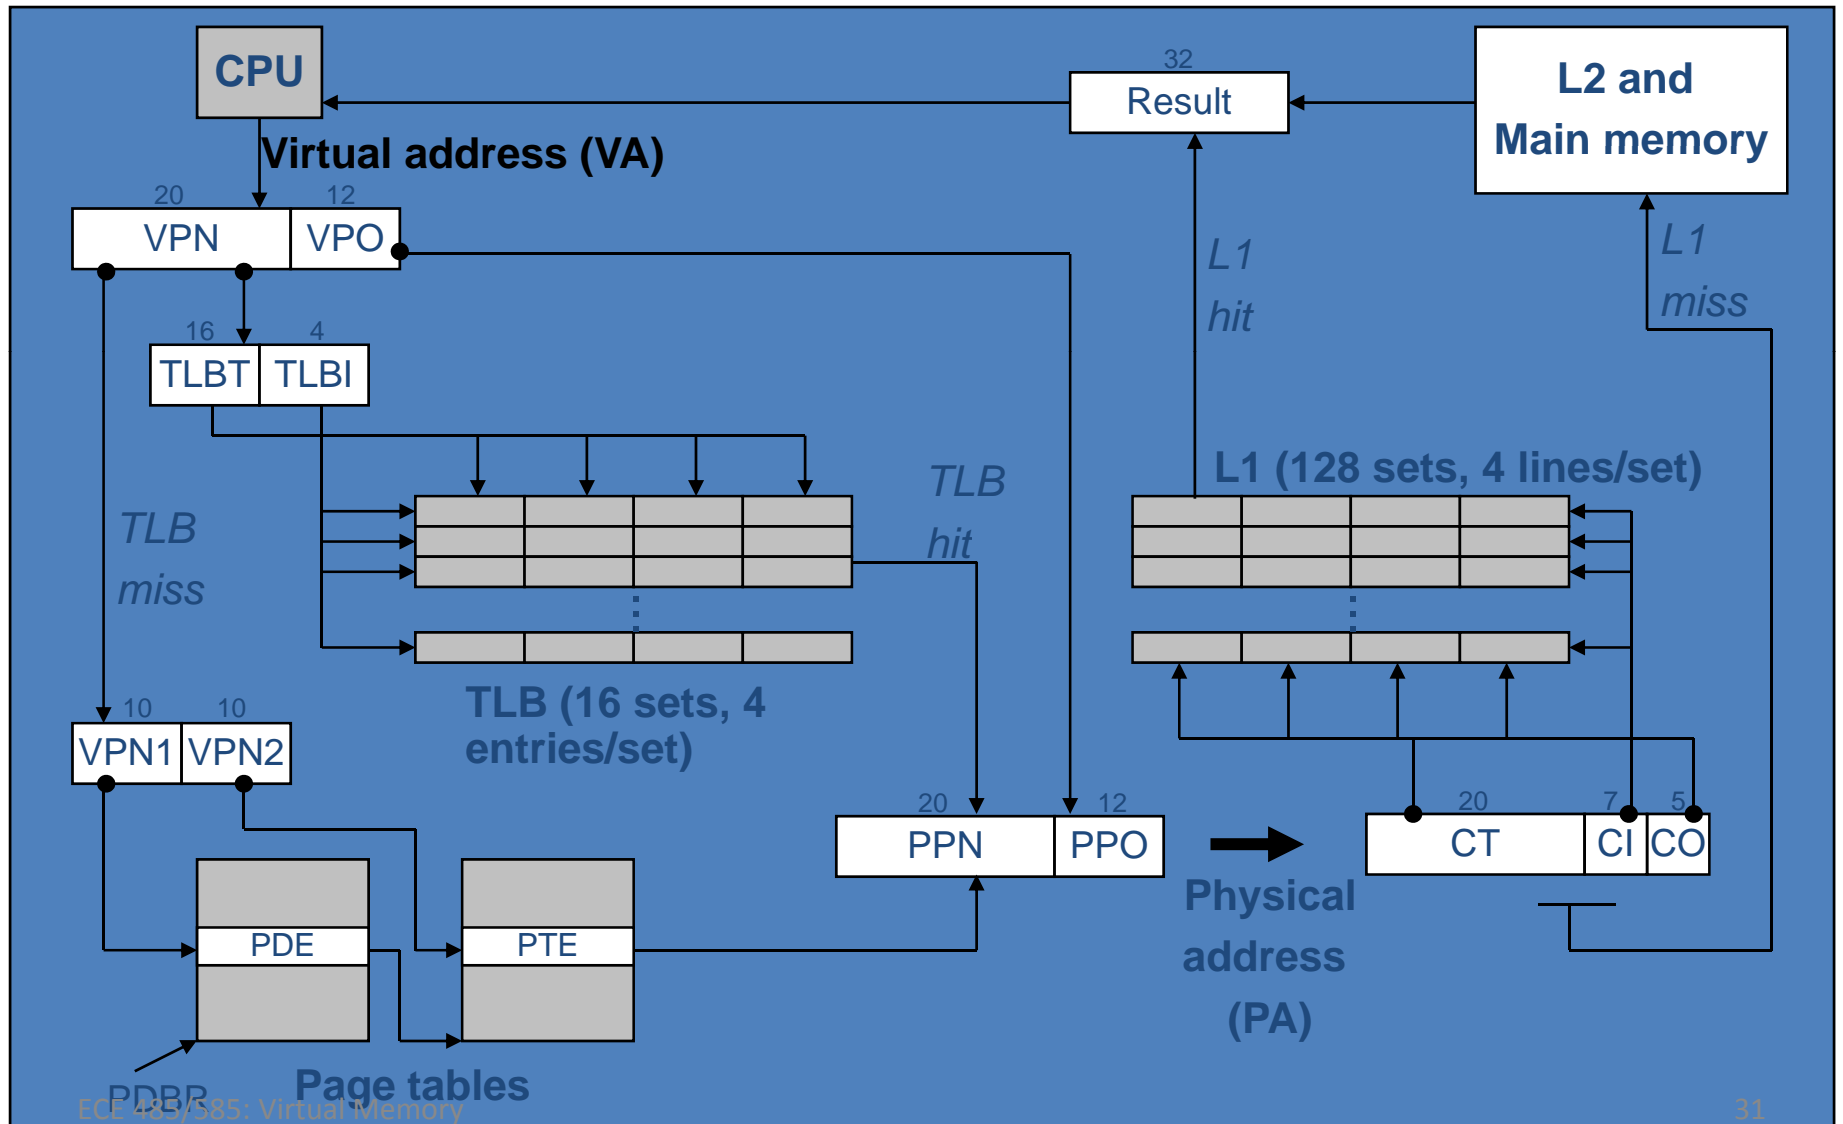
Lookup is in hardware, but maintenance (flush on context swap, replacement algorithm on miss, page table walking) can be handled in hardware (e.g. PowerPC, Intel x86) or software (Sparc, MIPS, ARM)

IA-32 terminology:  
Linear Address is Virtual Address after translation by segmentation hardware




LA = Linear Address  
PA = Physical Address

# Pentium Address Translation



# TLBs



↓ TLB (cache) lookup [often fully associative]

Virtual page number	C	U	D	W	G	OS	P	Page frame address

If entry present in TLB, don't need to walk the page table


If not present

Walk the page table


Evict an entry from TLB

Insert new mapping

Flush TLB (except for Global pages) whenever process switch!



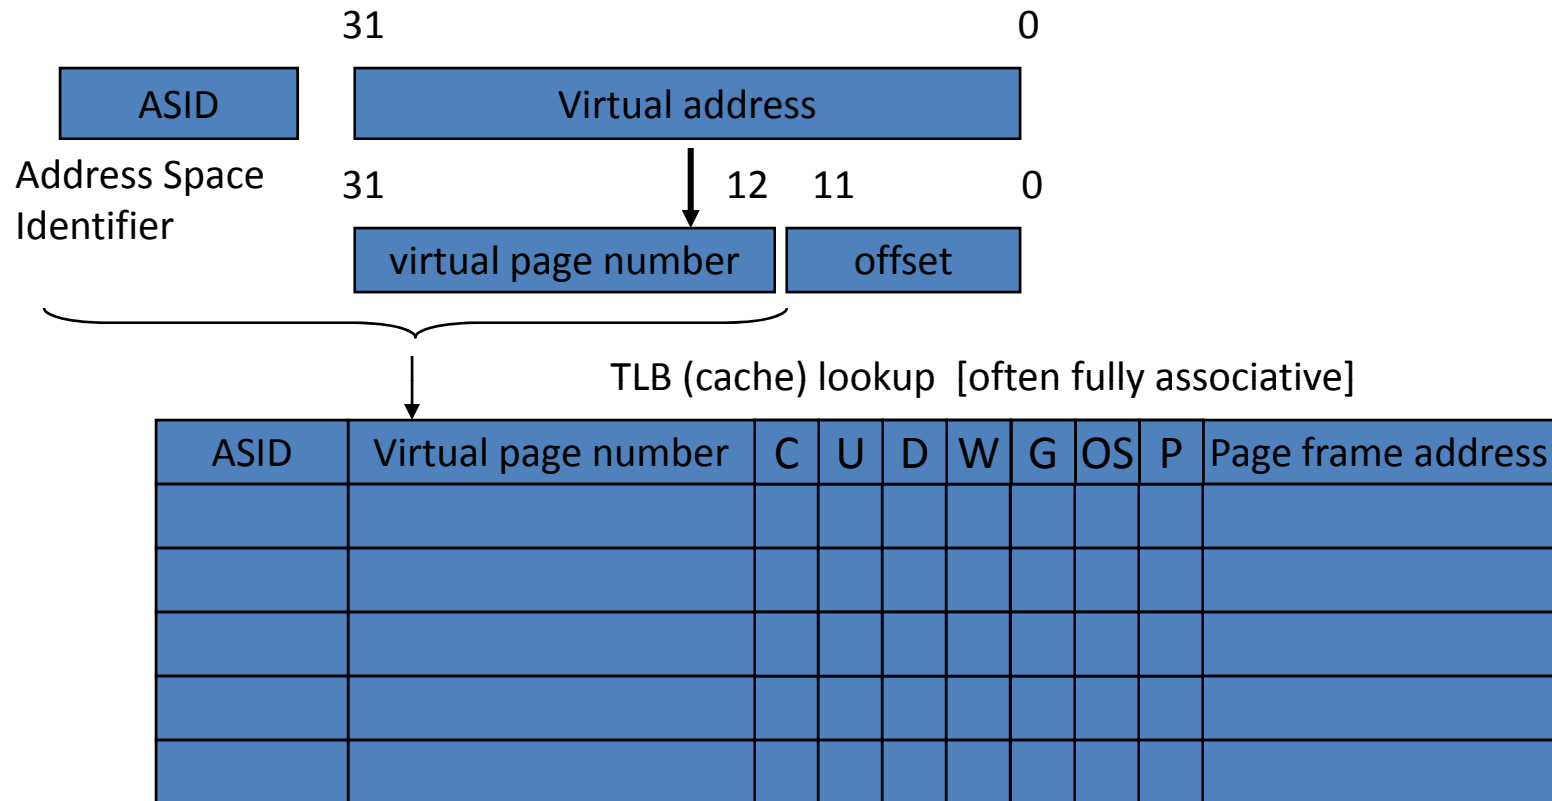
Page Table (Process A)



Page Table (Process B)<sub>32</sub>




# TLBs




By including ASID in tag we can allow PTEs from different processes to remain in the TLB to improve performance. No longer need to flush TLB on context switch. [Used in MIPS, Alpha, Sparc]. Global bit (G) avoids ASID comparison (OS, shared libraries)

New problem: Number of ASIDs < number of processes...

# Page Table Walking



# Top Down vs. Bottom Up Traversal



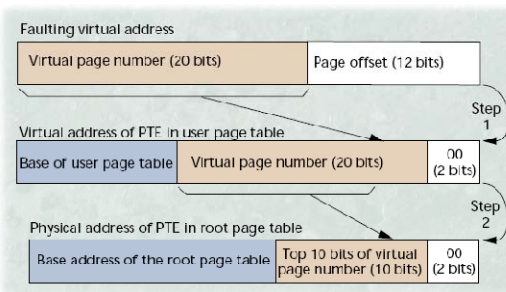

Top down traversal requires multiple memory references to obtain physical address.

If second level page table is stored in contiguous virtual memory addresses, a short cut is possible...


The index into the second level page table is the same as the virtual page number of the page it maps. We can create address from base of user page table and virtual page number and check TLB for this translation.

If found, load the PTE into the TLB. If not, obtain the physical address of the PTE entry in the first level page table for the second level page table's location

First used commercially in MIPS, later in Alpha



# Top Down vs. Bottom Up Traversal

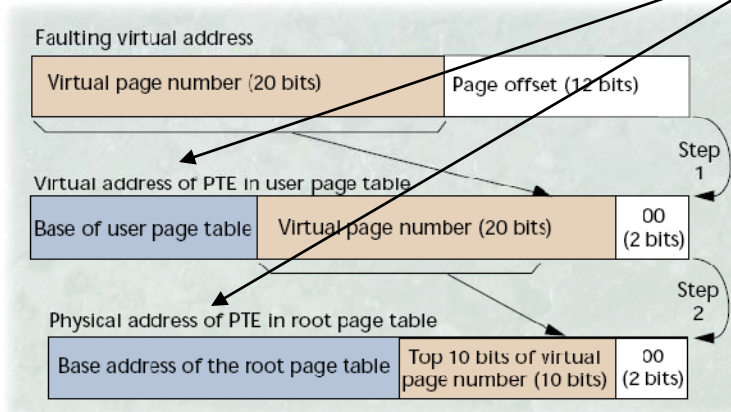


Physical address is 0xn<sup>20</sup>nnnn000 (4K)

20 bits

Virtual address is 0xnnC00000 (4M)

10 bits



(VPN, PPN) not in TLB

See if VPN of PTE is in TLB (if yes, use it to get physical address of page frame). LSBs = 00 because PTE is 4 bytes long.

No, get physical address of PTE from root page table

Note: page containing PTE and user page can both still result in page faults if not resident in memory. This process just loads the TLB for the VPN to PPN translation. The page may be non-resident in memory.


# Inverted Page Tables

Alternative: inverted page tables proportional to size of physical memory instead of virtual address space. [PowerPC, UltraSparc]

Hash anchor table not always used (omitted in PowerPC, PA-RISC) because it necessitates additional memory reference.

Downside: only maps pages which are currently resident in memory. Need another structure to map pages which are on disk so they can be paged in.

Can't accommodate different virtual addresses mapping to same physical address at same time. If memory is used for inter-process communication, may incur two page faults for each message.

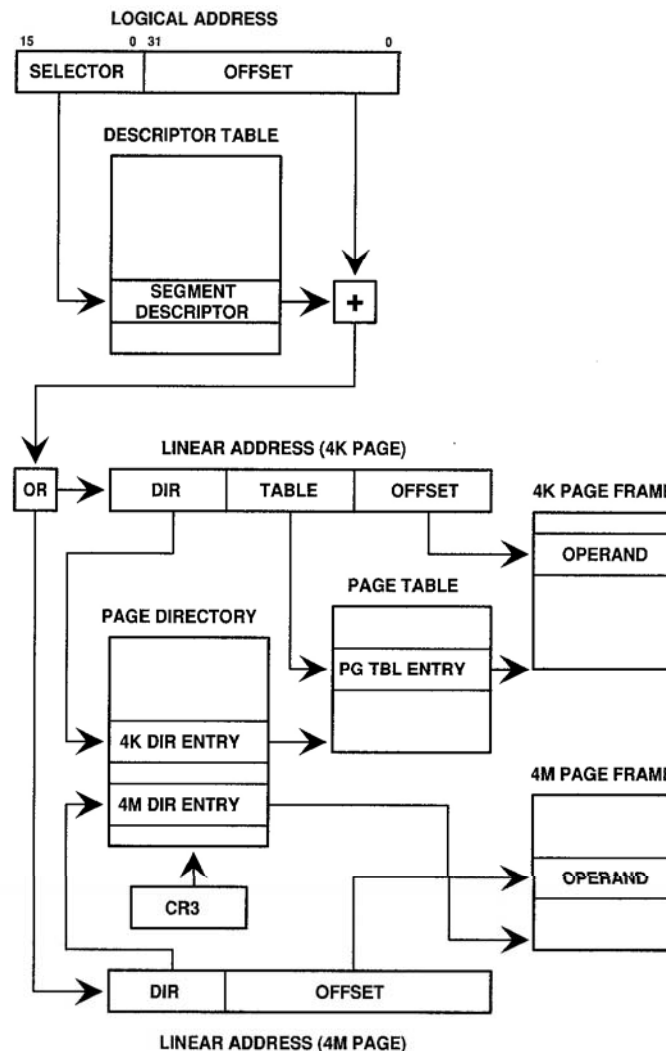


# Segmentation and Paging

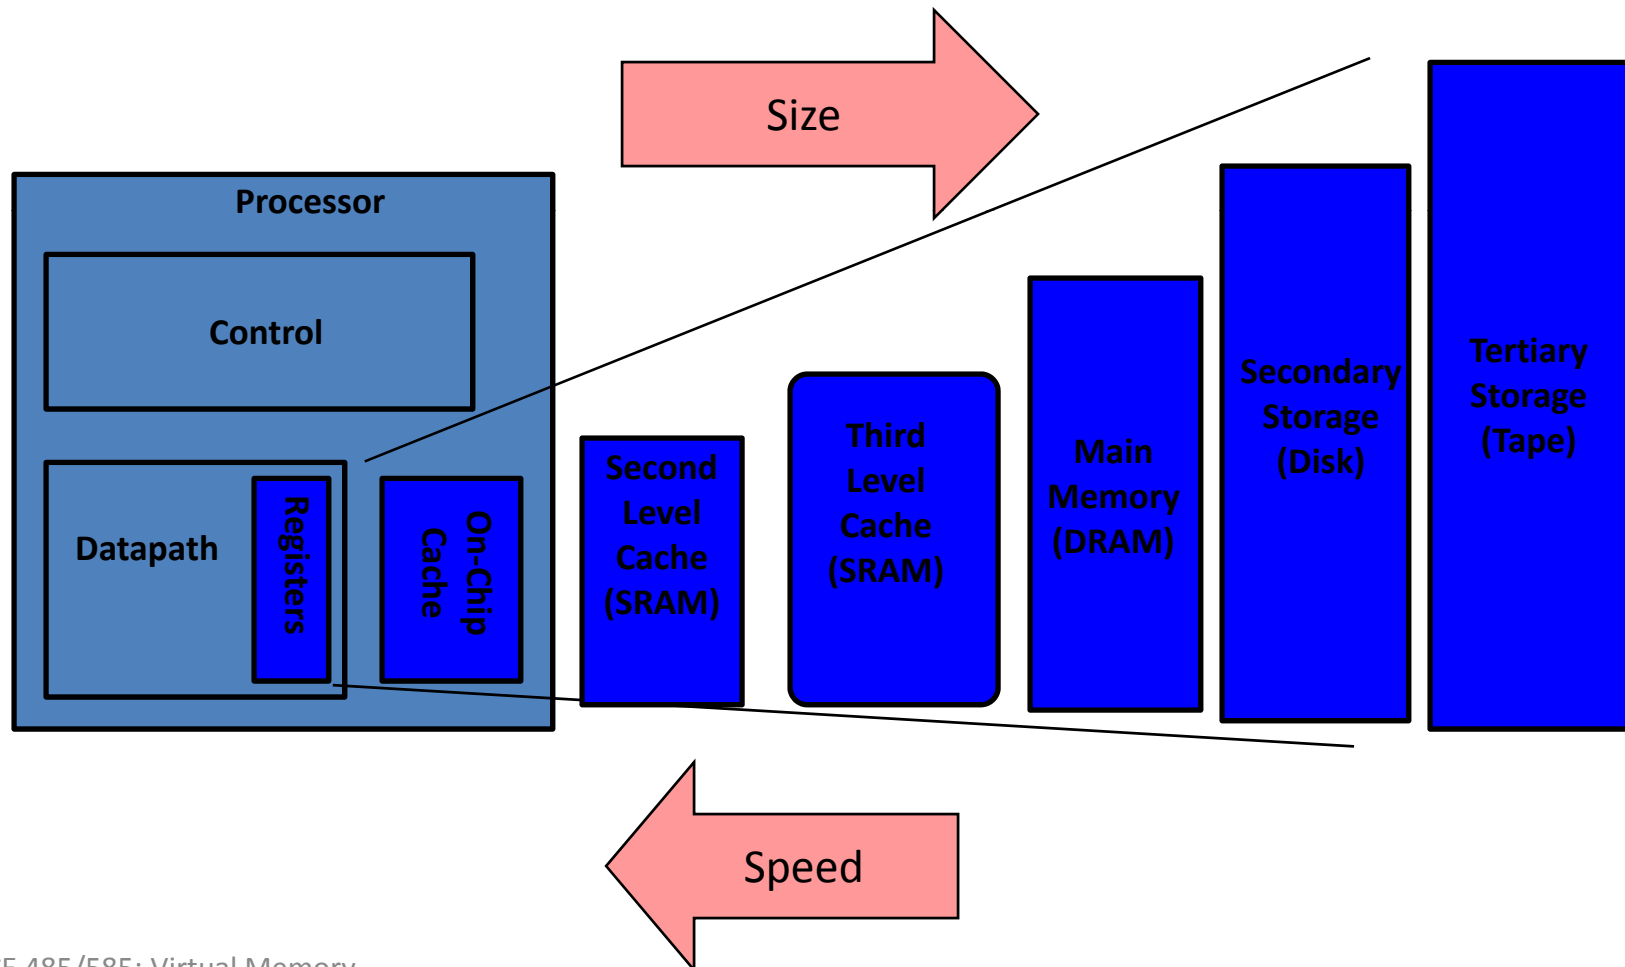
- Segmentation designed to provide very large virtual address space and protection
- Paging was designed to support efficient swapping between physical memory and disk
- However, paging provides adequate protection so virtually all modern operating systems (e.g. Windows NT, Unix) put all 32-bit code in single 4 GB segment. This is called the “flat address” model.
- Cannot disable segmentation on Pentium
- Effectively disable it by having single segment pointed to by all segment registers
- Use paging for everything else

# Combining Paging and Segmentation

Intel IA-32  
(also PowerPC)



# Memory Hierarchy of a Modern Computer System





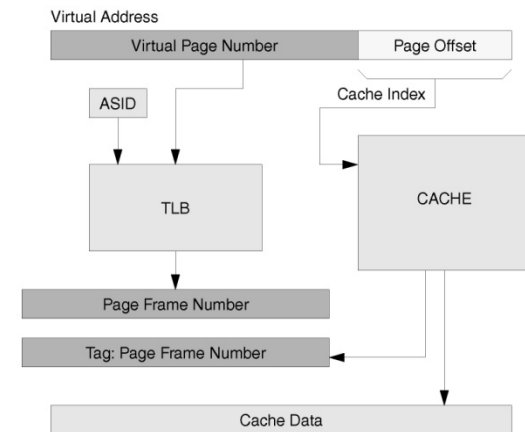
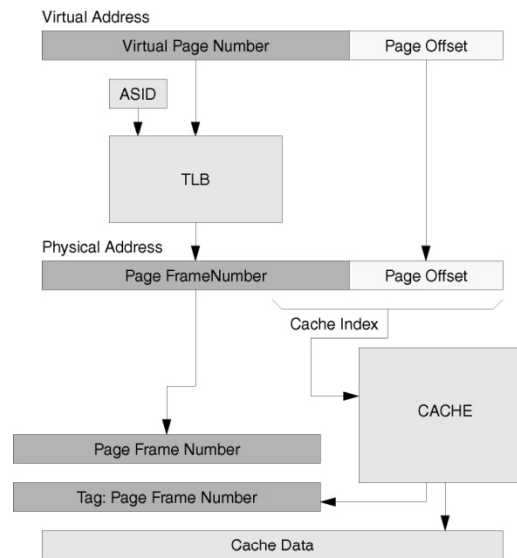

# Intel Pentium 4 3.2 GHz Server

Component	Access Speed (Time for data to be returned)
Registers	1 cycle = 0.3 nanoseconds
L1 Cache	3 cycles = 1 nanoseconds
L2 Cache	20 cycles = 7 nanoseconds
L3 Cache	40 cycles = 13 nanoseconds
Memory	300 cycles = 100 nanoseconds
Disk	15,000,000 cycles = 5,000,000 nanoseconds

# Caches and Virtual Memory

Parameter	L1 Cache	Virtual Memory
Block (page) size	16-128 bytes	4096-65,536 bytes
Hit time	1-3 clock cycles	50-150 clock cycles
Miss penalty	8-150 clock cycles	$10^6$ - $10^7$ clock cycles
Access time	6-130 clock cycles	$8 \times 10^5$ - $8 \times 10^6$ clock cycles
Transfer time	2-20 clock cycles	$2 \times 10^5$ - $2 \times 10^6$ clock cycles
Miss rate	0.1-10%	0.00001 – 0.001%
Address mapping	24-25 bit physical to 14-20 bit cache (index)	32-64 bit virtual to 25-45 bit physical

# Paging and Caching



## Advantage

Physically tagged → no synonym problem

## Disadvantage

Virtual address must be translated to physical

address first

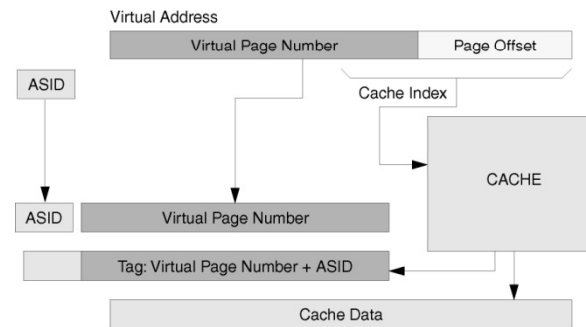

Must wait for TLB translation  
(or page table walk if miss)

## Solution

Size cache so that virtual page offset is same size as cache line index

Cache and TLB lookup in parallel

# Paging and Caching



## Advantage

- No TLB required

- Can use on cache miss

- Must store ASID in cache

- Different processes can use same virtual address for different physical locations

## Disadvantage

- Aliasing

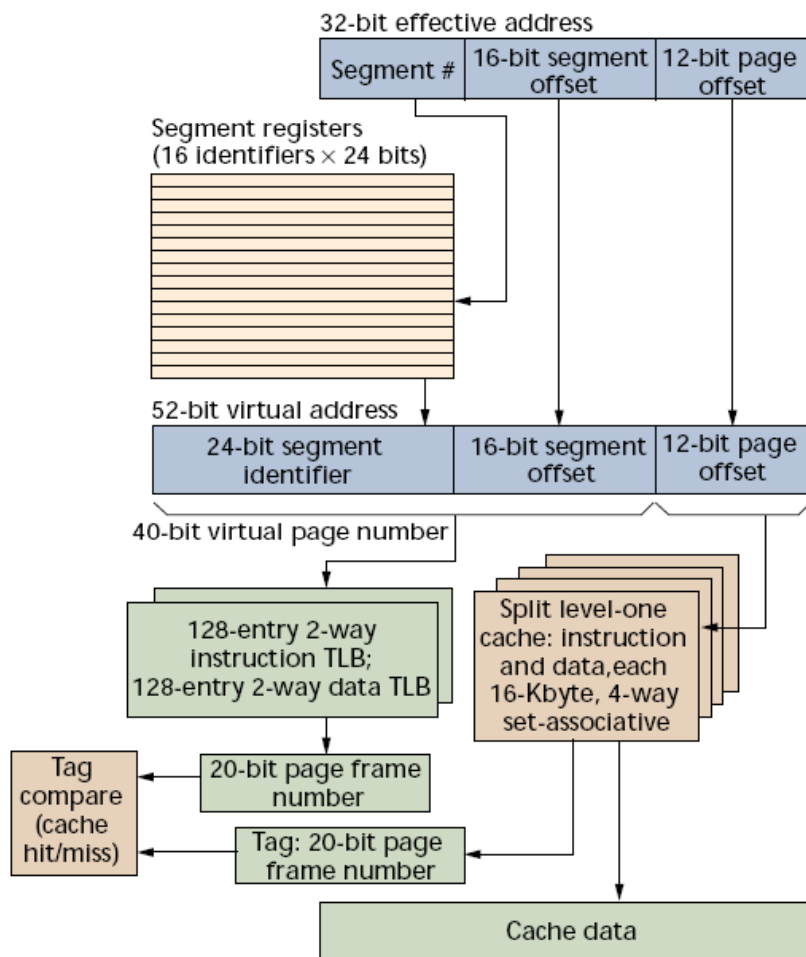
- Different physical addresses can map to same physical address at different cache locations!

## Solution

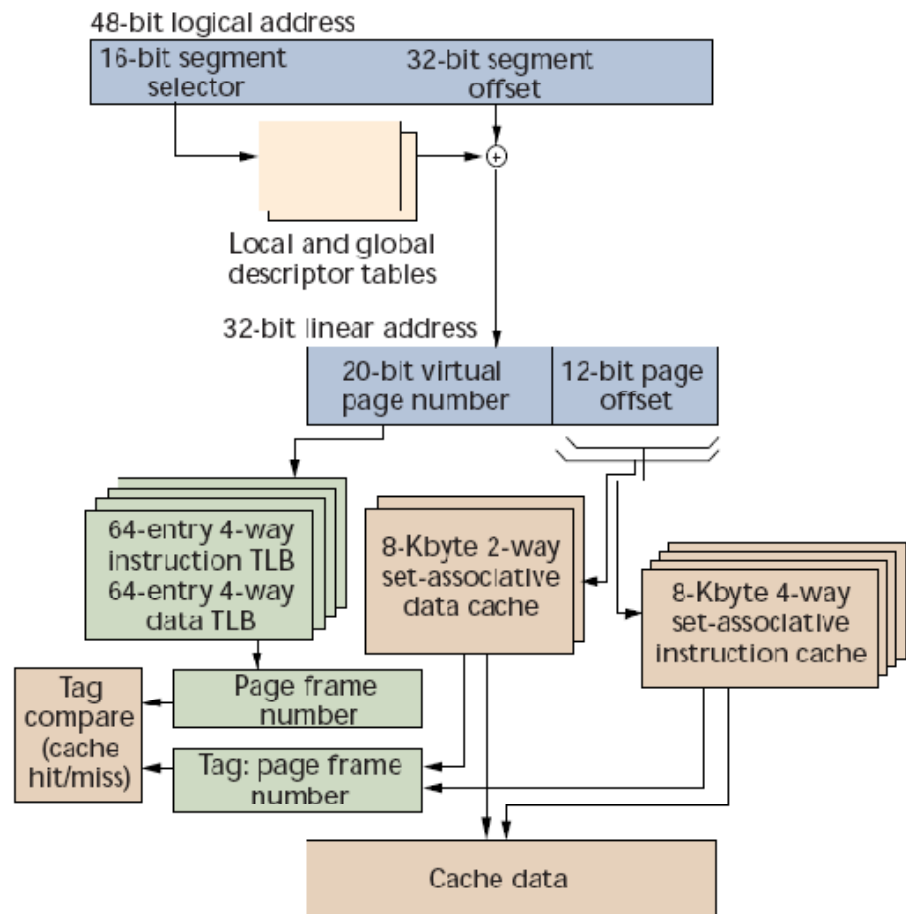
- Direct mapped cache only

Note: Still need some place to record physical address of cached on backing store for WB

# Some Examples

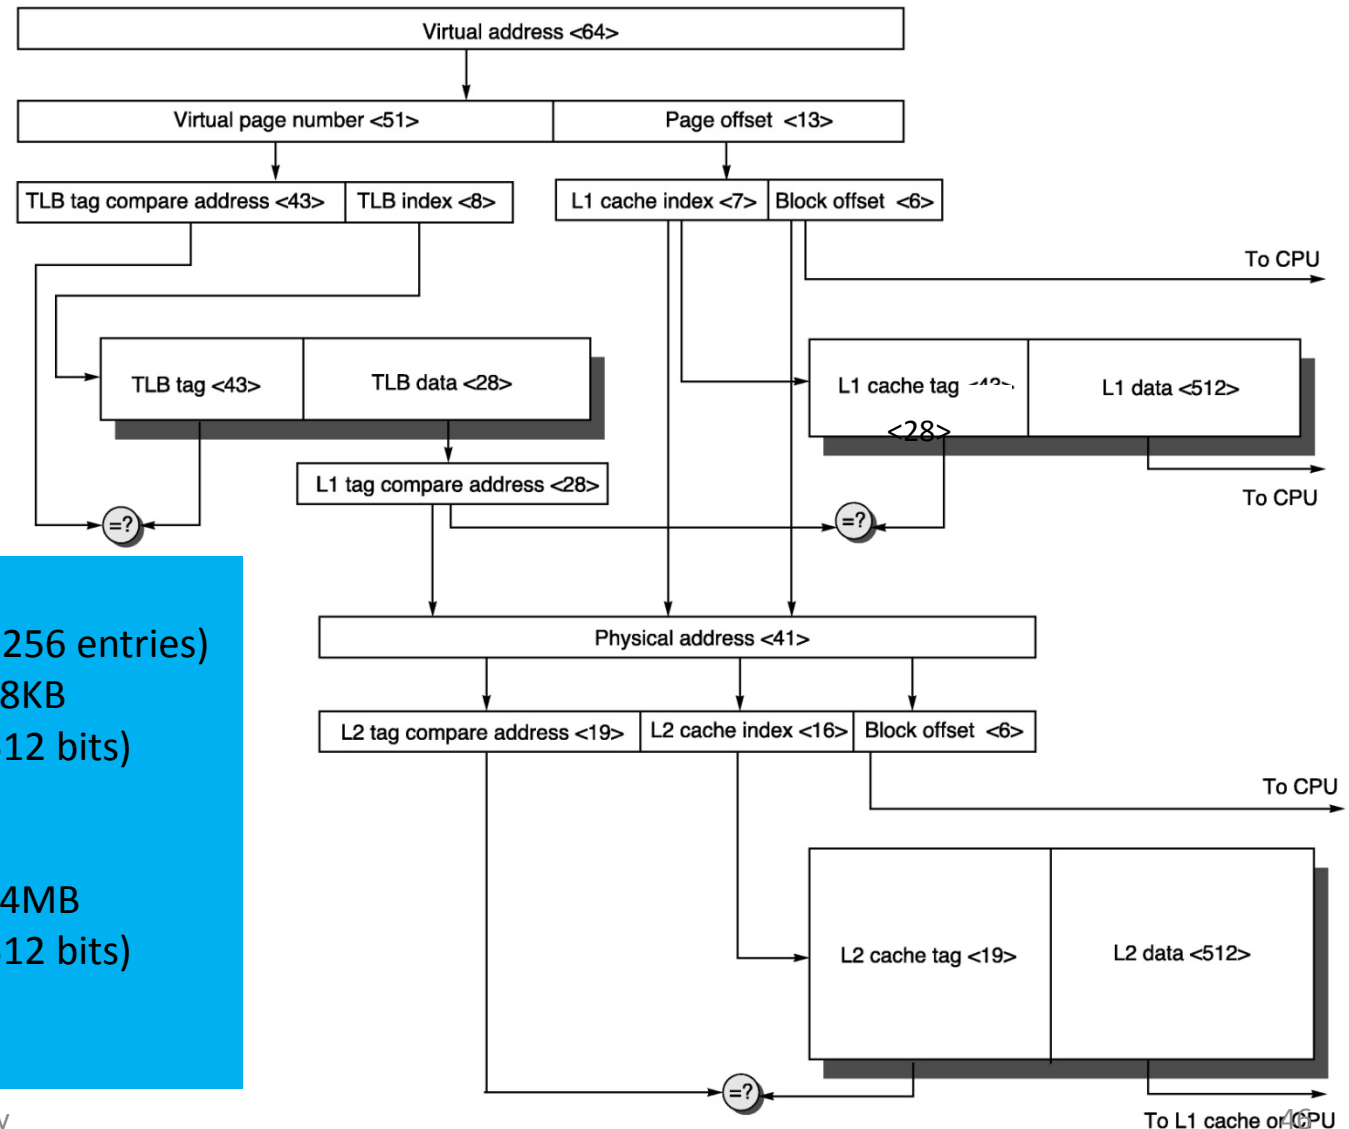


PowerPC 604



X86

# Paging and Caching



8K page size

TLB direct-mapped ( $2^8 = 256$  entries)

L1 cache direct-mapped 8KB

line size 64 ( $2^6$ ) bytes (512 bits)

7 bit index

$2^7 \times 2^6 = 2^{13} = 8K$

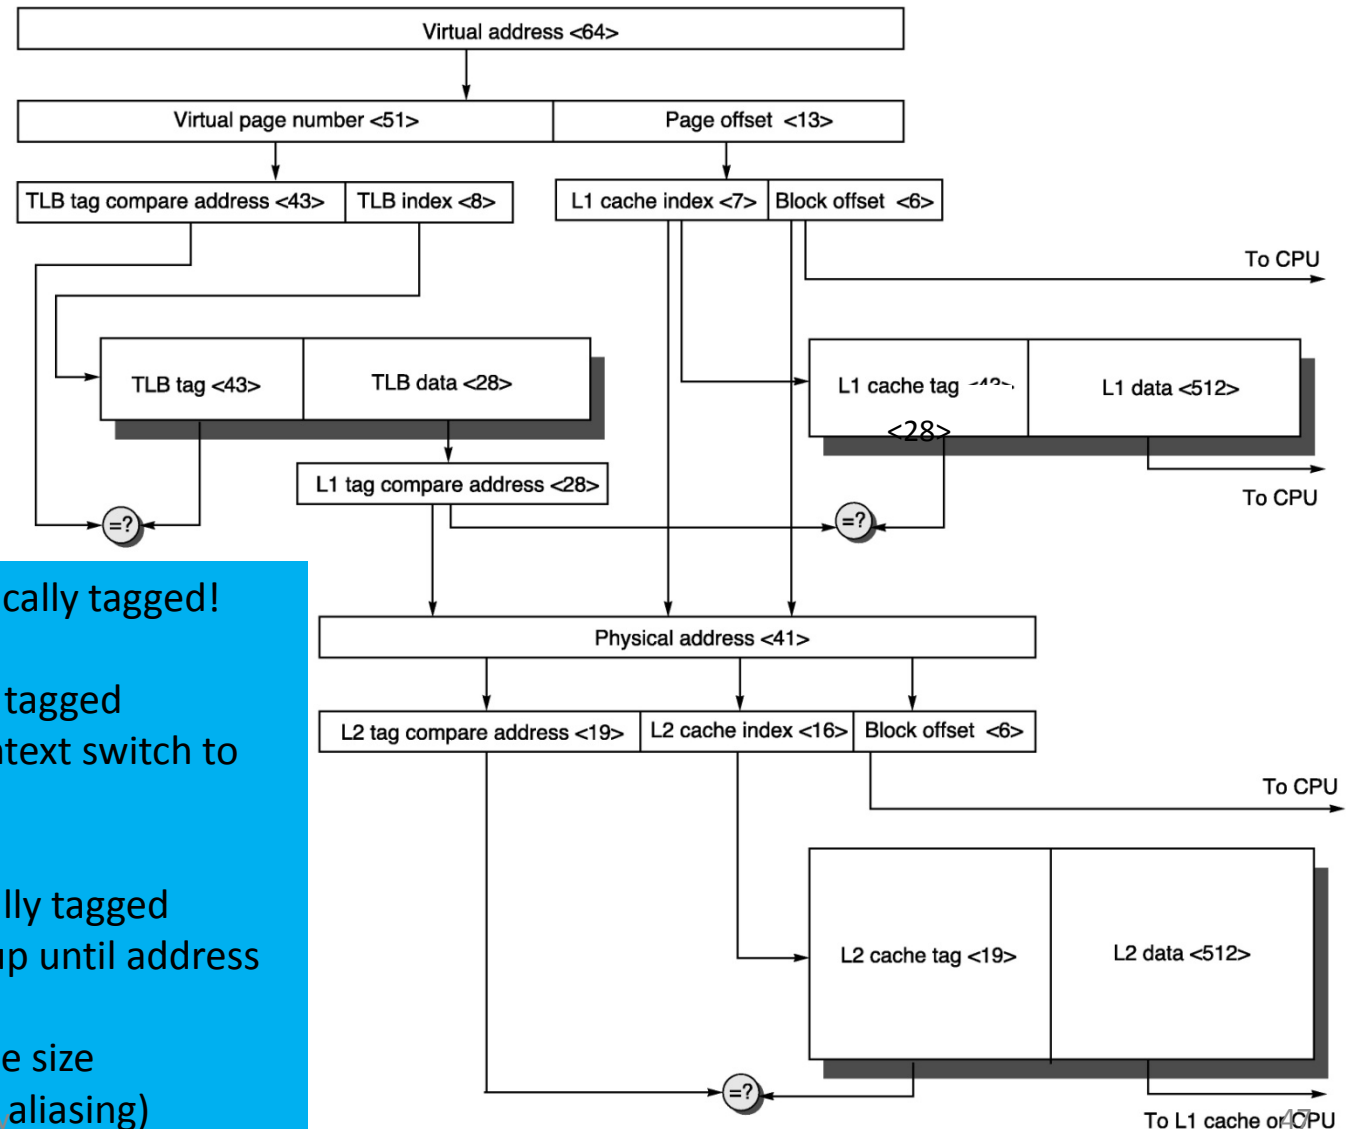
L2 cache direct-mapped 4MB

line size 64 ( $2^6$ ) bytes (512 bits)

16 bit index

$2^{16} \times 2^6 = 2^{22} = 4M$

# Paging and Caching



Virtually indexed but physically tagged!

Virtually indexed, virtually tagged

→ Must flush cache on context switch to avoid aliasing!

→ No need for TLB!

Physically indexed, physically tagged

→ Can't begin cache lookup until address translation complete

Simple if cache size == page size

(otherwise must deal with aliasing)

# Architectural Support for VM

Item	MIPS	Alpha	PowerPC	PA-RISC	UltraSPARC	IA-32
Address space protection	ASIDs	ASIDs	Segmentation	Multiple ASIDs	ASIDs	Segmentation
Shared memory	GLOBAL bit in TLB entry	GLOBAL bit in TLB entry	Segmentation	Multiple ASIDs; segmentation	Indirect specification of ASIDs	Segmentation
Large address spaces	64-bit addressing	64-bit addressing	52-/80-bit segmented addressing	96-bit segmented addressing	64-bit addressing	None
Fine-grained protection	In TLB entry	In TLB entry	In TLB entry	In TLB entry	In TLB entry	In TLB entry; per segment
Page table support	Software-managed TLB	Software-managed TLB	Hardware-managed TLB; inverted page table	Software-managed TLB	Software-managed TLB	Hardware-managed TLB/hierarchical page table
Superpages	Variable page size set in TLB entry: 4 Kbyte to 16 Mbyte, by 4	Groupings of 8, 64, 512 pages (set in TLB entry)	Block address translation: 128 Kbytes to 256 Mbytes, by 2	Variable page size set in TLB entry: 4 Kbytes to 64 Mbytes, by 4	Variable page size set in TLB entry: 8, 64, 512 Kbytes, and 4 Mbytes	Segmentation/variable page size set in TLB entry: 4 Kbytes or 4 Mbytes
ECE 485/585: Virtual Memory						48



# Intel Nehalem and AMD Opteron

Characteristic	Intel Nehalem	AMD Opteron X4 (Barcelona)
Virtual address	48 bits	48 bits
Physical address	44 bits	48 bits
Page size	4 KB, 2/4 MB	4 KB, 2/4 MB
TLB organization	<p>1 TLB for instructions and 1 TLB for data per core</p> <p>Both L1 TLBs are four-way set associative, LRU replacement</p> <p>The L2 TLB is four-way set associative, LRU replacement</p> <p>L1 I-TLB has 128 entries for small pages, 7 per thread for large pages</p> <p>L1 D-TLB has 64 entries for small pages, 32 for large pages</p> <p>The L2 TLB has 512 entries</p> <p>TLB misses handled in hardware</p>	<p>1 L1 TLB for instructions and 1 L1 TLB for data per core</p> <p>Both L1 TLBs fully associative, LRU replacement</p> <p>1 L2 TLB for instructions and 1 L2 TLB for data per core</p> <p>Both L2 TLBs are four-way set associative, round-robin</p> <p>Both L1 TLBs have 48 entries</p> <p>Both L2 TLBs have 512 entries</p> <p>TLB misses handled in hardware</p>

## AMD Opteron

48-bit virtual address  
40-bit physical address

### L1 TLB

40 entries, fully associative

### L2 TLB

4-way, 512 entries,  $\log_2(512/4) = 7$

### L1 data, instruction caches (separate)

64KB, 2-way, LRU replacement

64-byte cache line

$\log_2(64KB/(64 \times 2)) = 9$

### L2 unified cache

1M, 16-way, 64-byte cache line

pseudo-LRU replacement

### Victim buffer

8 entries

### L2 prefetcher

based on patterns of misses

### Memory controller

8 simultaneous data cache misses

2 simultaneous instruction misses

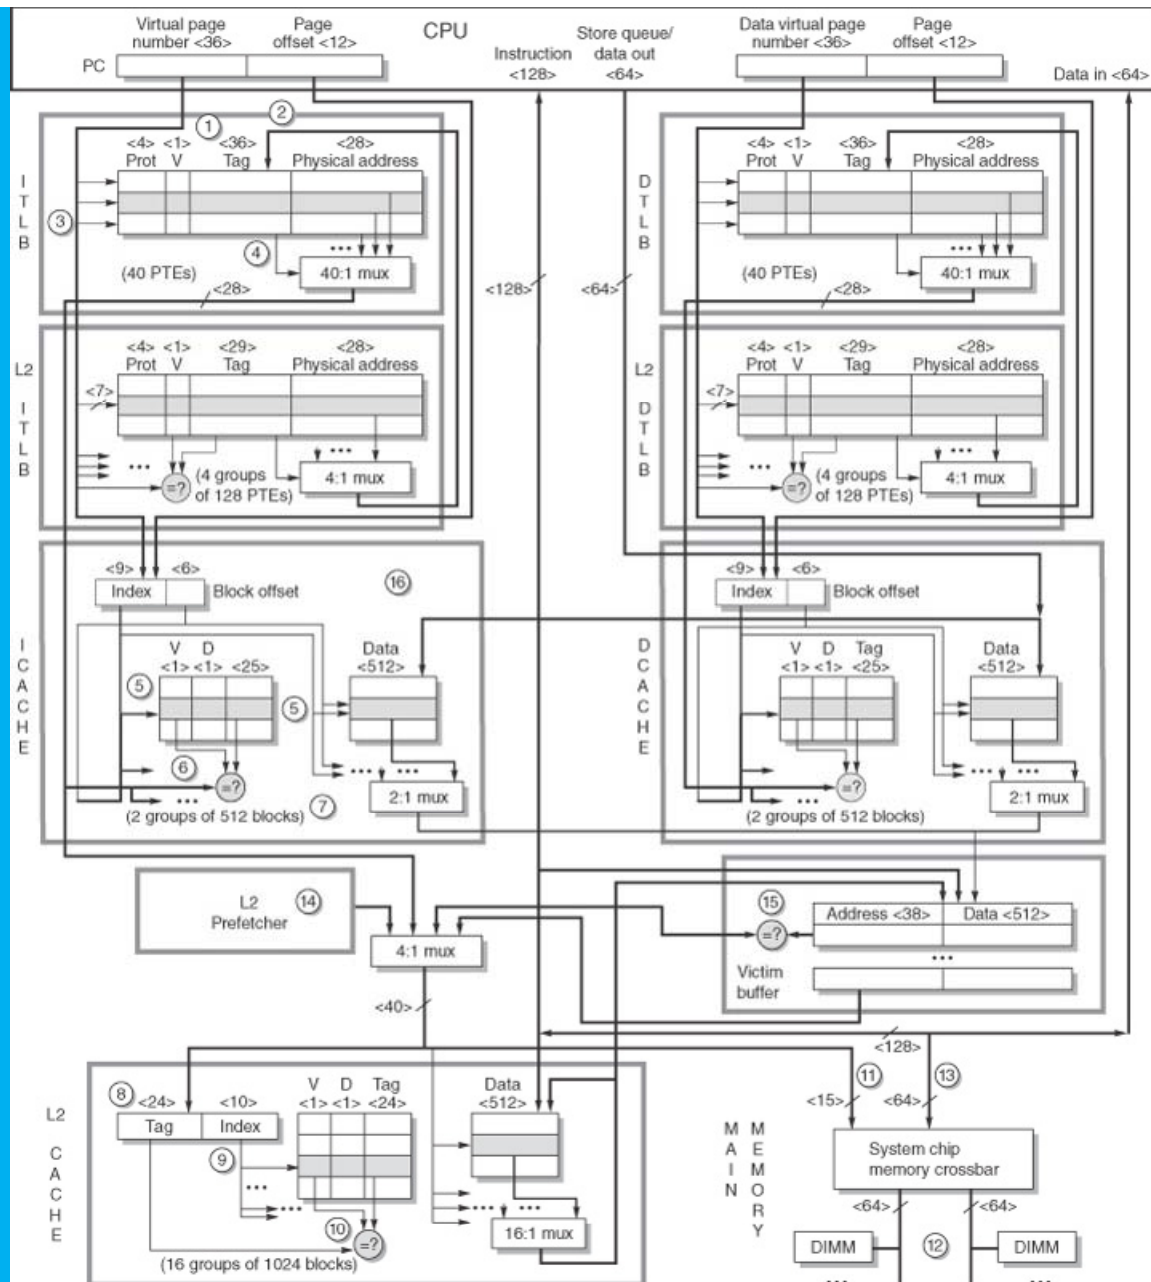
(hit under 10 miss)

2x 64-bit (+ECC) channels

### Physically indexed, virtually tagged

Page size not equal to cache size!

potential for aliasing – H/W checks!



# Benefits of Virtual Memory

- Programs have illusion of entire potential address space
  - All programs can start at address (e.g. 0) without being remapped by loader
- Don't need contiguous physical memory
  - Programs can co-exist, be swapped in/out with small granularity (no holes)
- Facilitates dynamic storage allocation (e.g. malloc)
  - Simply request more memory, pages allocated as needed
  - Physical memory required when used ("demand paged")
  - No need to copy/relocate code or data
  - Don't have issues of stacks growing into each other
- Protection
  - Hardware assisted range checking via page fault mechanism
    - Permissions on per page basis for sharing
    - Can detect referencing null pointer
    - Prevents accidental or malicious attempts to access memory not owned by process