

PORTLAND STATE UNIVERSITY

Microprocessor System Design

Project Report for Design and Simulation of a Split L1 Cache

Neerja Bawaskar (PSU ID: 912362484)

Rahul Wagh (PSU ID: 967287161)

Ameya Deswandikar (PSU ID: 970606576)

12/5/2012

Course Number: ECE585

Department: Electrical and Computer Engineering

Maseeh College of Engineering and Computer Science

Portland State University

CONTENTS

Project Specification.....	2
Design.....	3
Assumptions.....	3
Flowcharts.....	4
Read.....	4
Write.....	5
Reset.....	6
Invalidate.....	6
Testing.....	7
Test case design.....	7
Trace File 1.....	8
Trace File 2.....	9
Trace File 3.....	10
Trace File 4.....	11
Implementation.....	12
Choice of Language.....	12
Development environment.....	12
Class Diagram.....	13
File Dependencies.....	14
Error Handling.....	15
Maintainability.....	15
Repository.....	15
Coding Standards.....	16
Appendix.....	17
Source code.....	17
Trace file.....	17
Simulation results.....	17

PROJECT SPECIFICATION

Design and simulation of a split L1 cache backed by a unified L2 cache with MESI cache coherence protocol in a multiprocessor system.

Specifications:

Parameter	L1 Data Cache	L1 Instruction Cache
Associativity	Four-way	Two-way
Number of sets	16K	16K
Line size	64 bytes	64 bytes

Both caches implement an LRU replacement policy and have a shared L2 cache.

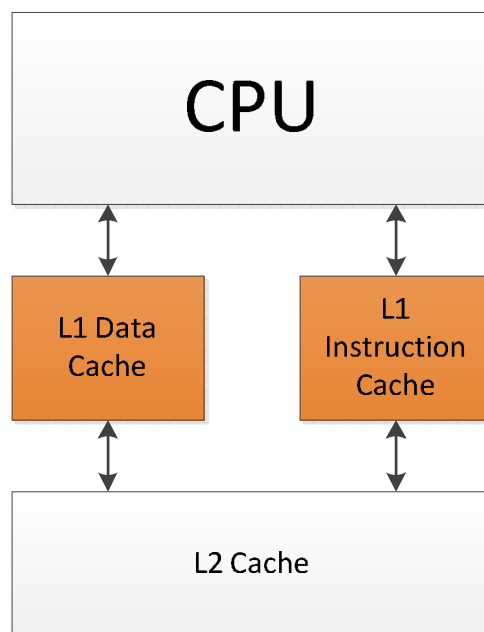


Figure 1: L1 Cache external interfaces

DESIGN

The L1 cache is supported by a unified L2 cache. Since this is a multiprocessor system with MESI cache coherence mechanism, we maintain inclusivity with L2 to simplify the L1 cache design.

In order to maintain inclusivity, we implemented a write-through with write allocate policy on both L1 Data and Instruction caches. This implies that we need not implement the MESI protocol in the L1 cache. The cache coherence mechanism is handled in the lower level cache (L2), freeing up the L1 cache to serve the CPU. Design and simulation of the L2 cache is out of scope for this project. We tested the L1-L2 interface using a L2 stub: a minimal implementation of the cache interface specified below.

The figure below represents the public interfaces provided by the L1 Data and Instruction caches.

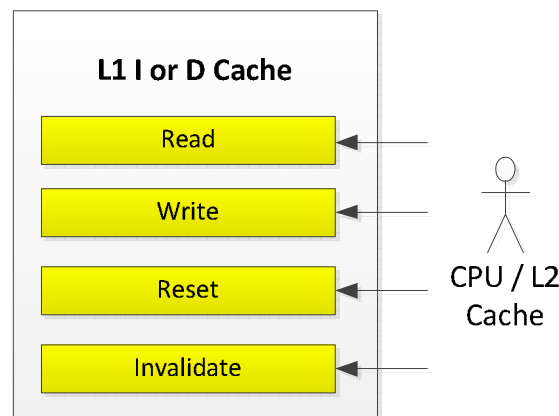


Figure 2: Public Interfaces

Interface	Description
Read	Read data from the cache
Write	Write data to the cache
Reset	Clears the Valid and Tag bits and clears the data in the cache
Invalidate	Clears the Valid bit of the cache line

ASSUMPTIONS

- The last level in the memory hierarchy has to be Write Back in order to satisfy the given specification of write back with respect to memory.
- The memory is byte addressable.
- The maximum size of address is 32 bits.

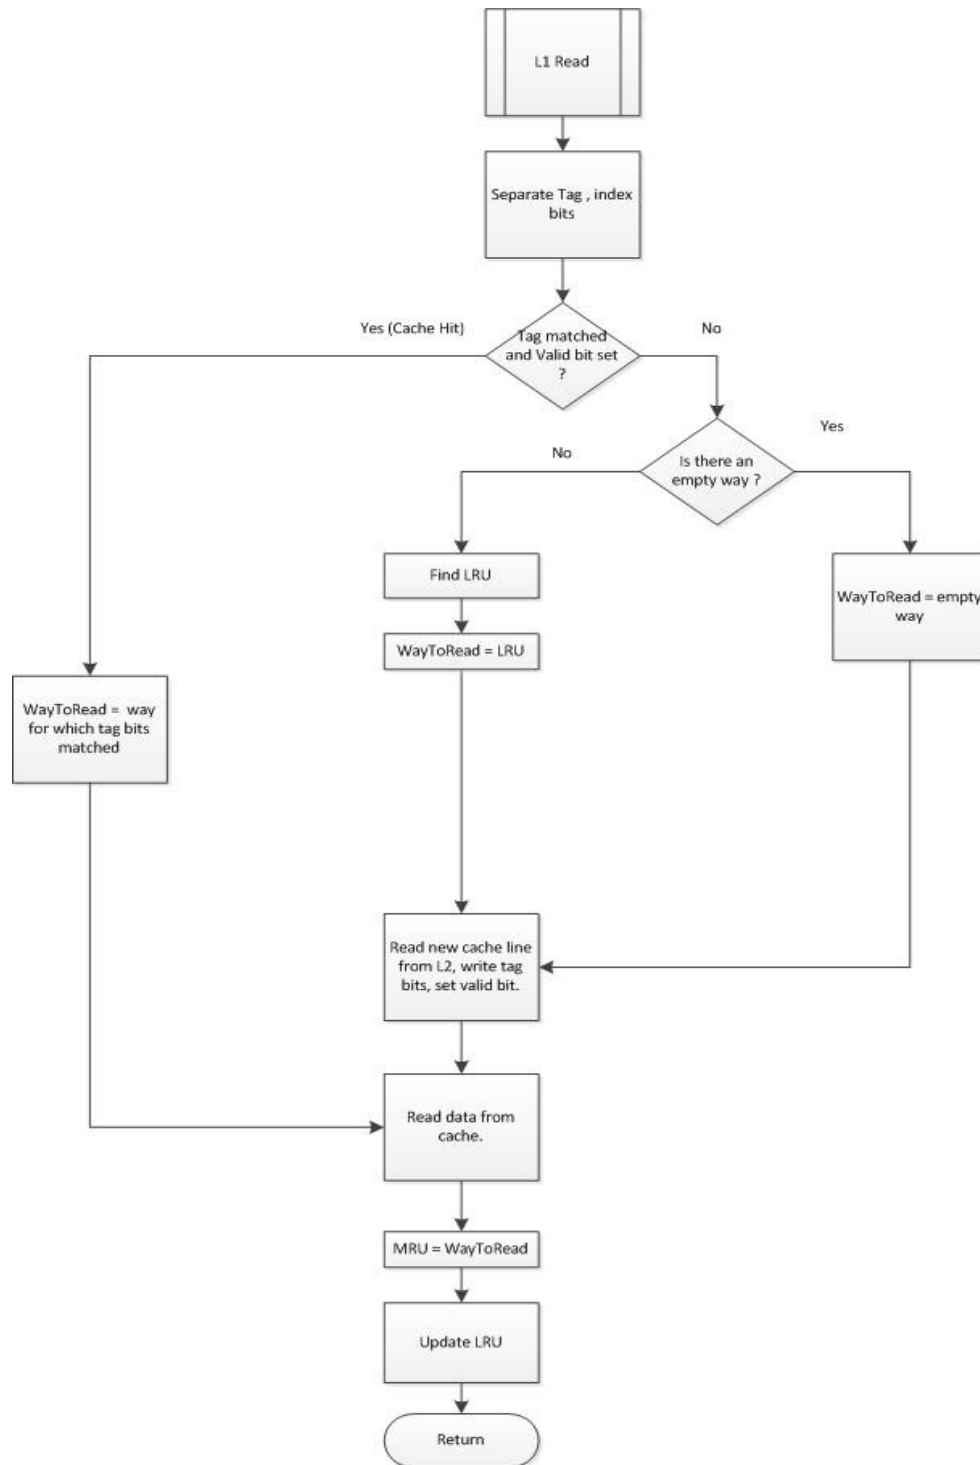
FLOWCHARTS**READ**

Figure 3: Flowchart for READ

WRITE

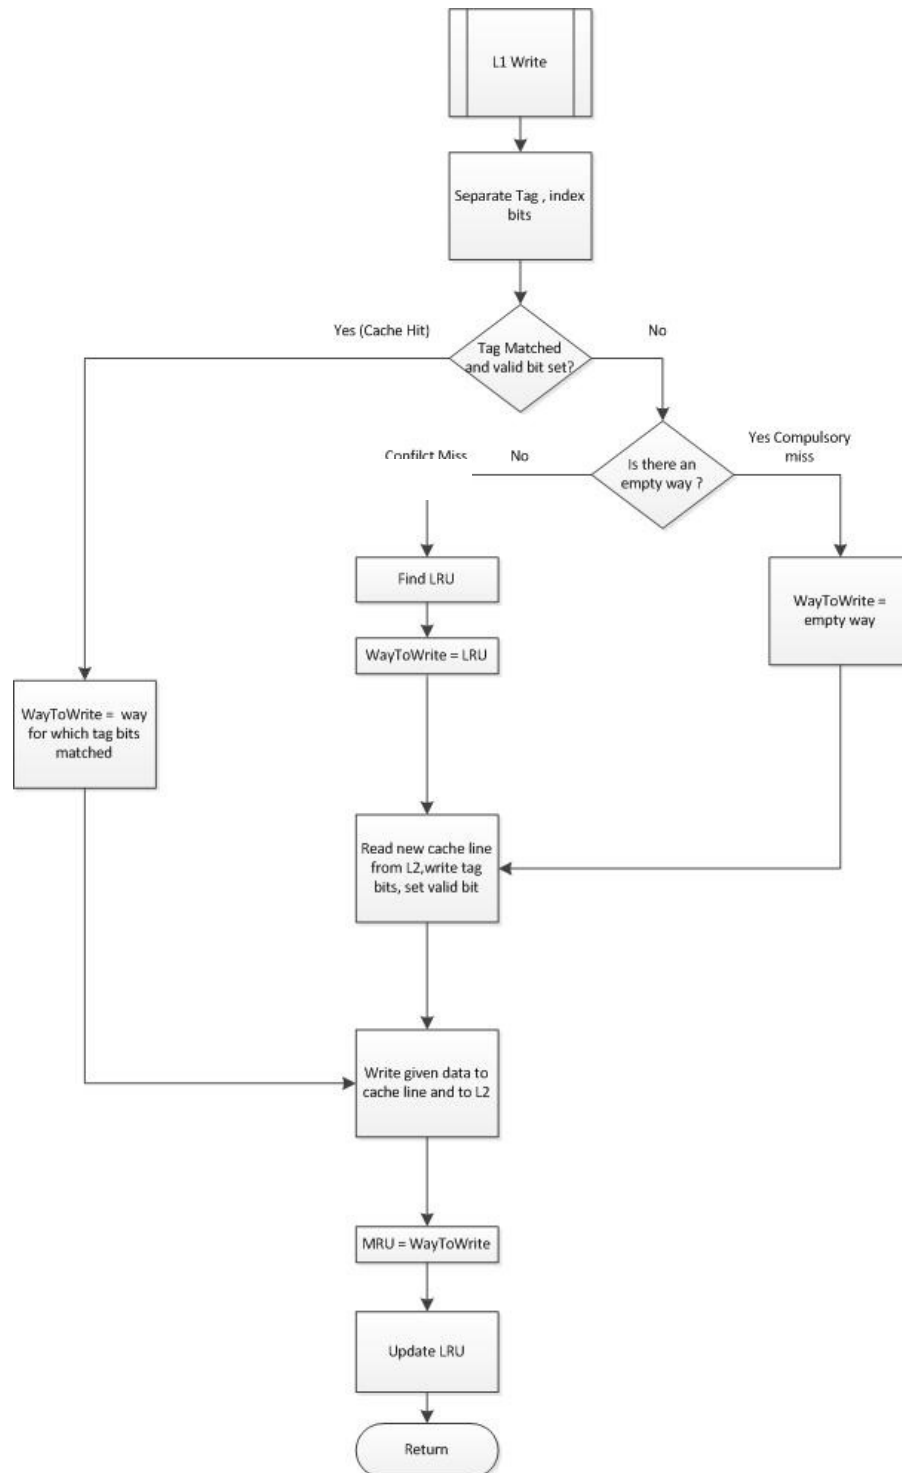


Figure 4: Flowchart for WRITE

RESET

The Reset command clears the Valid bits and the Tag bits of each line in the cache. The LRU bits are set to a default value of 3 (i.e. least recently used). The data in each line of every set is also cleared.

INVALIDATE

The Invalidate command makes the cache line invalid by clearing the Valid bit of the line. The reasons for invalidate could be due to an eviction in L2 cache (due to inclusivity) or due to the cache coherence protocol implementation (MESI). In case of a hit to a modified line, L2 cache will send an Invalidate to the L1.

TESTING

The trace file has the following format:

n address

where n is

0	Read data request from L1-D cache
1	Write data request to L1-D cache
2	Instruction Fetch (Read request from L1-I cache)
3	Invalidate command from L2
8	Clear the cache and reset all states
9	Print contents and state of the cache

TEST CASE DESIGN

We designed test cases to validate the working of our cache in the following scenarios. The test cases were implemented using four different trace files.

1. Write hits to each line in the set. (L1 Data cache only).
2. Reads hits to each line in the set. (For both L1 Data and Instruction cache)
3. Eviction cases of all lines in the set. (For both L1 Data and Instruction cache)
4. Testing an Invalidate command to show how it affects a hit.
5. Testing a Cache hit for different bytes in the same line.
6. Cache hits to the set with at least one empty line
7. Reset and Print commands
8. Print the cache statistics when the entire trace file is read
9. Testing for invalid command inputs i.e. other than 0,1,2,3,8 and 9
10. Testing for addresses greater than 32-bits
11. Testing for whitespace and blank lines

TRACE FILE 1

1. Only writes to L1-D
2. Last set accessed
3. Invalidate command
4. Eviction cases for all lines in the set using LRU replacement policy
5. Cache write hit for different bytes in the same line
6. Write hits to all lines in the set

Table 1: Trace file 1 entries

Command	Address	Tag Bits (12)	Index Bits (14)	Offset Bits (6)	Hit/Miss	Comments
1	0xFFFC0	0x000	0x3FFF	0x0	Miss	Way = 0
1	0xE0FFFC1	0xE00	0x3FFF	0x1	Miss	Way = 1
3	0xFFFC8	0x000	0x3FFF	0x8	---	Invalidate Way 0
1	0xFFFC7	0x000	0x3FFF	0x7	Miss	Way = 0
1	0xD0FFFC0	0xD00	0x3FFF	0x0	Miss	Way = 2
1	0xC0FFFF7	0xC00	0x3FFF	0x37	Miss	Way = 3
1	0xB0FFFC7	0xB00	0x3FFF	0x7	Miss, Evict	LRU = 1
1	0xA0FFFC5	0xA00	0x3FFF	0x5	Miss, Evict	LRU = 0
1	0x90FFFD3	0x900	0x3FFF	0x13	Miss, Evict	LRU = 2
1	0x80FFFE2	0x800	0x3FFF	0x22	Miss, Evict	LRU = 3
1	0xB0FFFC7	0xB00	0x3FFF	0x7	Hit	Way = 1
1	0xA0FFFC5	0xA00	0x3FFF	0x5	Hit	Way = 0
1	0x90FFFD3	0x900	0x3FFF	0x13	Hit	Way = 2
1	0x80FFFE2	0x800	0x3FFF	0x22	Hit	Way = 3

For L1 Data Cache:

No. of reads = 0

No. of writes = 13

No. of misses = 9

No. of hits = 4

Hit Ratio = 0.3076

TRACE FILE 2

1. Only reads to both L1-D and L1-I
2. Hits to a cache line in the set which has at least one empty line.
3. Cache read hit for different bytes in the same line
4. Read hits to all lines in the set for both caches

Table 2: Trace file 2 entries

Command	Address	Tag Bits (12)	Index Bits (14)	Offset Bits (6)	Hit/Miss	Comments
0	0xF000FFF0	0xF00	0x3FF	0x30	Miss	Way=0
0	0x E000FFF0	0xE00	0x3FF	0x30	Miss	Way=1
0	0x F000FFF7	0xF00	0x3FF	0x37	Hit	Way=0
0	0x D000FFF0	0xD00	0x3FF	0x30	Miss	Way=2
0	0x C000FFF0	0xC00	0x3FF	0x30	Miss	Way=3
0	0x B000FFF0	0xB00	0x3FF	0x30	Miss, Evict	LRU=1
0	0x A000FFF0	0xA00	0x3FF	0x30	Miss, Evict	LRU=0
0	0x 9000FFF0	0x900	0x3FF	0x30	Miss, Evict	LRU=2
0	0x 8000FFF0	0x800	0x3FF	0x30	Miss, Evict	LRU=3
0	0x B000FFF0	0xB00	0x3FF	0x30	Hit	Way=1
0	0x A000FFF0	0xA00	0x3FF	0x30	Hit	Way=0
0	0x 9000FFF0	0x900	0x3FF	0x30	Hit	Way=2
0	0x 8000FFF0	0x800	0x3FF	0x30	Hit	Way=3
2	0x F000FFF0	0xF00	0x3FF	0x30	Miss	Way=0
2	0x F000FFF1	0xF00	0x3FF	0x31	Hit	Way=0
3	0x F000FFF1	0xF00	0x3FF	0x31	---	Invalidate Way=0
2	0x F000FFF1	0xF00	0x3FF	0x31	Miss	Way=0
2	0x E000FFF0	0xE00	0x3FF	0x30	Miss	Way=1
2	0x D000FFF0	0xD00	0x3FF	0x30	Miss, Evict	LRU=0
2	0x C000FFF0	0xC00	0x3FF	0x30	Miss, Evict	LRU=1
2	0x D000FFF0	0xD00	0x3FF	0x30	Hit	Way=0
2	0x C000FFF0	0xC00	0x3FF	0x30	Hit	Way =1

Table 3: Trace file 2 statistics

Parameter	L1 Data Cache	L1 Instruction Cache
Number of reads	13	8
Number of writes	0	0
Number of misses	8	5
Number of hits	5	3
Hit Ratio	0.3846	0.375

TRACE FILE 3

1. Both Reads and Write to L1-D cache tested
2. Consecutive Read, Write, Invalidate and Read to the same line tested
3. Print command tested
4. Reset command tested
5. Printing statistics at the end of line

Table 4: Trace file 3 entries

Command	Address	Tag Bits (12)	Index Bits (14)	Offset Bits (6)	Hit/Miss	Comments
1	0x4003C0	0x400	0xF	0x0	Miss	Way=0
0	0xF003C2	0xF00	0xF	0x0	Miss	Way=1
1	0xF003C4	0xF00	0xF	0x4	Hit	Way=1
3	0xF003C2	0xF00	0xF	0x0	---	Invalidate Way1
0	0xF003C2	0xF00	0xF	0x0	Miss	Way=1
9					Print	Print stats
8					Reset	Clear
0	0xF003C2	0xF00	0xF	0x0	Miss	Way=0

For L1 Data Cache:

Parameter	After '9' command	After end of file
Number of reads	2	1
Number of writes	2	0
Number of misses	3	1
Number of hits	1	0
Hit Ratio	0.25	0

TRACE FILE 4

1. Presence of an invalid command (other than 0,1,2,3,8 and 9)
2. Address greater than 32-bit - Extra bits starting from MSB truncated
3. Additional White spaces in the trace file

Table 5: Trace file 4 entries

Command	Address	Tag Bits (12)	Index Bits (14)	Offset Bits (6)	Hit/Miss	Comments
7	0x 4003C0	0x400	0xF	0x0	---	Bad Command
0	0x F00F003C2	0xF00	0xF	0x0	Miss	Address greater than 32 bits
2	0x F000FFFO	0xF00	0x3FF	0x30	Miss	White space tested

IMPLEMENTATION

This section deals with the implementation details of the simulator.

CHOICE OF LANGUAGE

We chose C++ to implement the simulator as opposed to Verilog or C. Key driver behind this was the implementation of two caches: L1 Instruction and L1 Data and their interface to another cache. We modeled a generic cache using an abstract class. This provides a consistent interface for all caches. The L1 data and instruction caches implement this interface as a write through cache. Design of the L2 stub also became easier, since it minimally implements this common interface.

DEVELOPMENT ENVIRONMENT

We used Microsoft Visual Studio C++ 2012 environment to design, develop and test our simulator. The project was built as an executable on the Windows Platform.

CLASS DIAGRAM

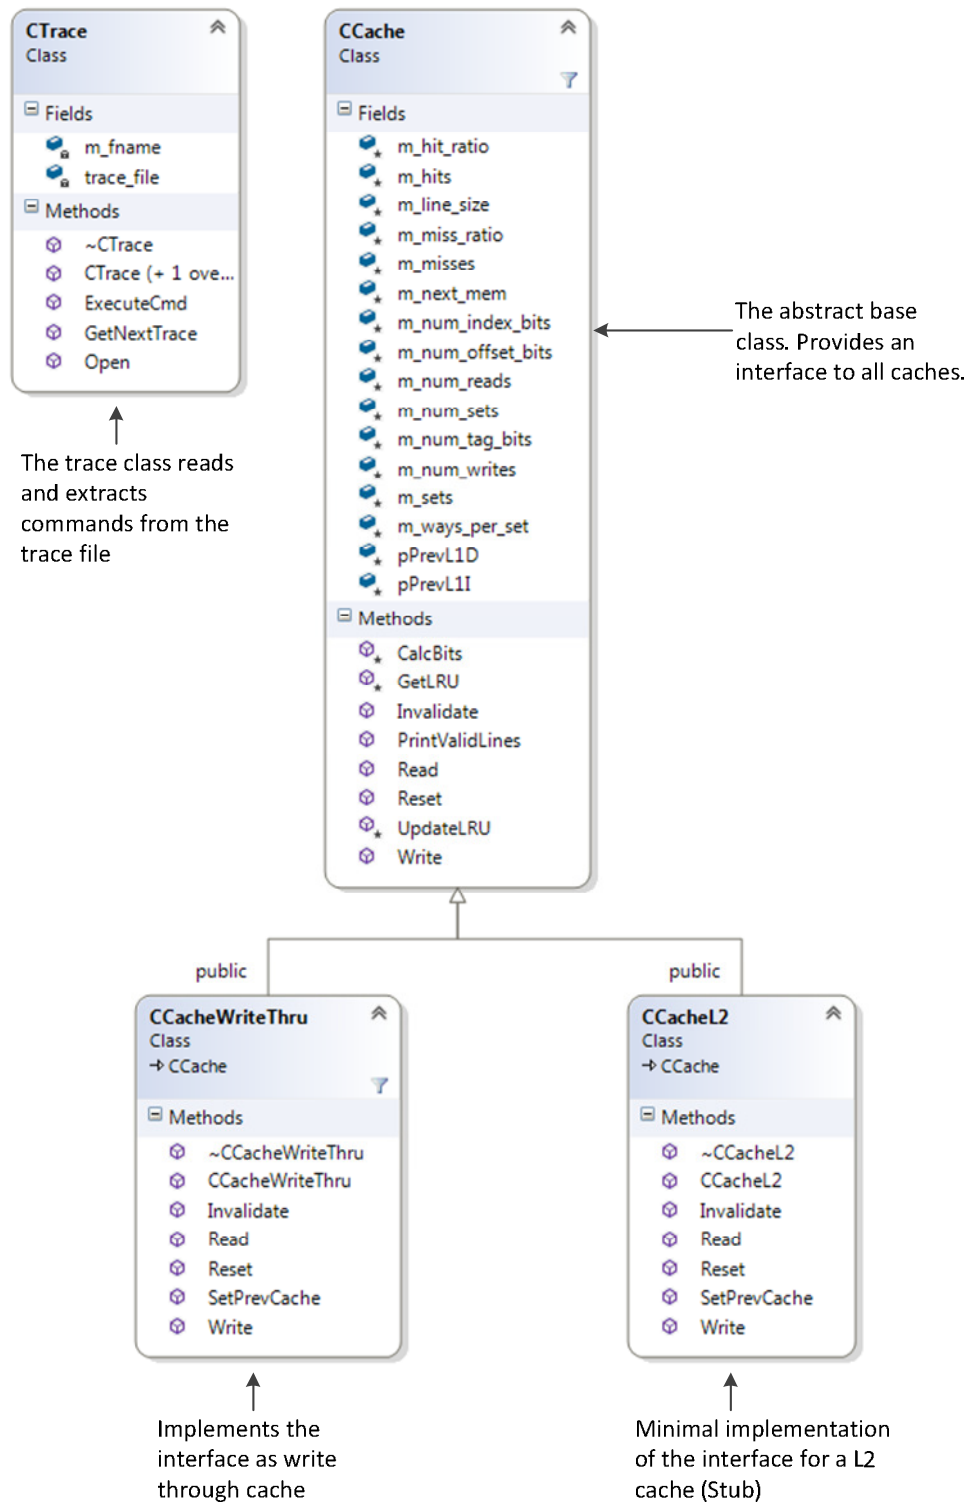


Figure 5: Class Diagram

FILE DEPENDENCIES

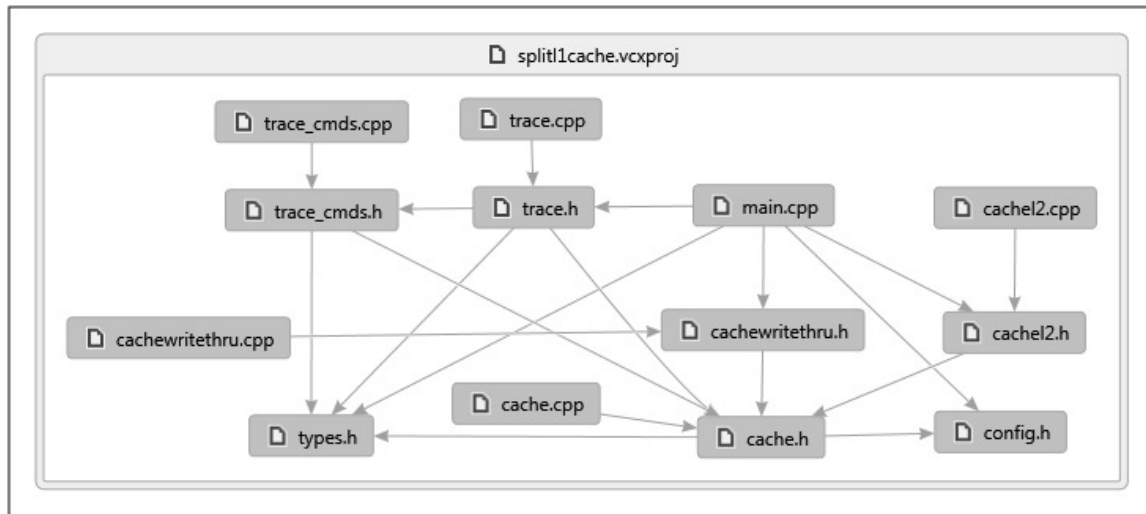


Figure 6 : File Dependency Graph

Table 6: File Organization

File	Module	Description
Types.h	Global	Contains type definitions for the entire project.
Config.h	Global	Configuration parameters such as Line size, number of sets etc
Cache.h	Cache	Public interface for class Cache
Cache.cpp	Cache	Implementation of class Cache
CacheWriteThru.h	CacheWriteThru	Public interface for class CacheWriteThru
CacheWriteThru.cpp	CacheWriteThru	Implementation of class CacheWriteThru
CacheL2.h	CacheL2	Public interface for class CacheL2
CacheL2.cpp	CacheL2	Implementation of CacheL2
Trace.h	Trace	Public interface for class Trace
Trace.cpp	Trace	Implementation of class Trace
Trace_cmds.h	Trace	Trace command function headers
Trace_cmds.cpp	Trace	Trace command function implementations
Main.cpp	Global	The integration of all modules

ERROR HANDLING

1. Arguments were checked before proceeding in any function.
2. Exceptions handling used to catch memory allocation failures for the new operator.
3. All pointers checked for NULL before dereferencing
4. All new lines / empty lines in trace file skipped
5. Invalid command line arguments checked
6. Invalid commands in the trace file tagged as bad commands and skipped.
7. Memory leaks checked using the `_CrtDumpMemoryLeaks()` function of Microsoft Windows C runtime.

MAINTAINABILITY

The trace commands can be easily extended. To add new commands, the user should add his command to the enumeration in `types.h`, assign the proper number to the enumeration (command number). Then the user can define the implementation of his command in `trace_cmds.h` / `.cpp`. Finally, he should add the enumeration and function pointer in the array `trace_commands` in `trace.h`. This provides a modular way for adding commands without use of complex switch case based implementations. This also results in minimal code changes.

The class design provides an abstract base class which serves as an interface for any future caches to implement. This is generic enough for all kinds of caches. They could make use of this to interact with each other in a consistent fashion.

REPOSITORY

We used Subversion for source control. The repository is hosted on Google Code and provides up to 5 GB space for hosting projects. All previous revisions can be easily obtained via the repository. It also contains an issue tracking system for managing defects.

CODING STANDARDS

Coding standards are used to

1. Organize code better
2. Eliminate possible hazards
3. Maintain consistency
4. Improve readability
5. Ease of maintenance
6. Achieve Portability

We used the C++ standard recommended by Google Inc. ([The Google C++ Style Guide](#)) Some features which picked up are:

1. Use of custom primitive data types: UINT8, UINT16 etc. (Portability)
2. One header file and cpp file per class. (Modularity)
3. Consistent naming schemes for Classes, variables, enums (Consistency, Readability)
4. Use of literal as first comparison argument e.g. (0 == my_var) (Avoid possible pitfalls)
5. Use of const references whenever possible (Optimal implementation)

APPENDIX

SOURCE CODE

TRACE FILE

SIMULATION RESULTS